

Master Engineering Systems

Major Project Report

Developing a Scalable Computer-Vision Pipeline for Human–Interaction Element Detection in Hospital Environments

Pavel Petrovski

HAN University of Applied Sciences

Supervisors: Siddharth Ajaykumar, Jan Benders, Ben Pyman

March 31, 2026



Master Engineering Systems – Major Project

Cyber-Physical Systems

HAN University of Applied Sciences

Arnhem

Preface

This report presents the work carried out as part of the Master thesis project in Cyber-Physical Systems track of the Master Engineering Systems program at HAN University of Applied Sciences. The project was conducted in collaboration with Ambyon and HAN Automotive Research.

The objective of the thesis is to design and evaluate a scalable computer-vision pipeline for service-robot interaction-point perception in hospital environments. All work described in this report was performed by the author.

The project repository is available at [GitLab](#), and experiment tracking together with fine-tuning curves is available at [Weights & Biases](#).

Disclaimer

This Project Report was written independently by the author as part of the Master of Engineering Systems program at HAN University of Applied Sciences.

During the preparation of this report, the generative artificial intelligence tools *ChatGPT* and *Codex* were used selectively to support the writing and development process. ChatGPT was used to assist with improving sentence structure, refining academic language, and suggesting minor text rephrasing. Codex was used for code clarity, bug fixing, and code sanity checks.

All technical content, system design decisions, experimental implementation, training procedures, evaluation methods, and results presented in this report are the original work of the author. The use of the AI tool did not replace independent analysis, interpretation of scientific literature, or critical reasoning.

The use of AI complies with the *MES Policy on the Use of Generative AI*, ensuring informed, transparent, ethical, and responsible application.

Summary

This report presents work carried out within the CareBOTS projects at HAN University of Applied Sciences in collaboration with Ambyon. The project addresses the need for reliable robot perception of elevator buttons in hospital environments, where autonomous navigation across floors depends on accurate detection, interpretation, and spatial localization of panel interfaces.

This work develops a perception workflow for autonomous robot interaction with elevator panels under embedded-hardware constraints. The approach combines instance segmentation for button localization, optical character recognition for semantic interpretation, and iterative model improvement to support practical deployment in hospital-like environments.

The results show that the final system can combine button segmentation, text recognition, and depth-based target estimation in a modular embedded perception pipeline running on NVIDIA Jetson hardware. The project therefore delivers a practical foundation for elevator-button interaction on the CareBOTS platform and for future model improvement through deployment-time data collection and retraining.

Nomenclature

ADC	Active Data Collector.
AI	Artificial Intelligence.
AP	Average Precision.
APA	American Psychological Association.
CNN	Convolutional Neural Network.
COCO	Common Objects in Context.
CPU	Central Processing Unit.
CUDA	Compute Unified Device Architecture.
CV	Computer Vision.
CVAT	Computer Vision Annotation Tool.
DL	Deep Learning.
FP16	Half-precision floating point format.
FPS	Frames Per Second.
GPU	Graphics Processing Unit.
IEEE	Institute of Electrical and Electronics Engineers.
IoU	Intersection over Union.
JSON	JavaScript Object Notation.
mAP	Mean Average Precision.
Mask R-CNN	Mask Region-Based Convolutional Neural Network.
ML	Machine Learning.
NVIDIA Jetson	Embedded GPU platform for edge AI.
OCR	Optical Character Recognition.
ONNX	Open Neural Network Exchange.
PaddleOCR	Paddle Optical Character Recognition.
RGB	Red–Green–Blue.
RGB-D	Red–Green–Blue plus Depth.
RNN	Recurrent Neural Network.
ROS	Robot Operating System.
SLAM	Simultaneous Localization and Mapping.
SOTA	State of the Art.
TBD	To Be Determined.
TensorRT	TensorRT.
VRAM	Video Random Access Memory.
W&B	Weights & Biases.
WSL2	Windows Subsystem for Linux 2.
XML	Extensible Markup Language.
YAML	YAML Ain't Markup Language.
YOLO	You Only Look Once.

Contents

Preface	i
Disclaimer	ii
Summary	iii
1 Introduction	1
1.1 Background	1
1.2 Problem Definition	2
1.3 Project Objective	2
1.4 Research Question(s)	2
1.5 Outline of the report	3
2 Literature Review	4
2.1 Perception of human-machine interaction elements	4
2.2 Formulating interaction-point perception as a vision problem	4
2.3 Public datasets and annotation strategies.	6
2.4 Model design principles, evaluation metrics, and semantic interpretation	6
2.5 Embedded deployment constraints in robotic perception	7
2.6 Scalability, robustness, and reproducibility.	8
2.7 Summary and identified research gap	8
3 Methodology	9
3.1 Overview of the proposed perception workflow	9
3.2 Dataset selection and preparation	9
3.3 Annotation strategy and workflow	11
3.4 Dataset quality assessment	12
3.5 Model architectures and training strategy	13
3.5.1 Experimental hardware and software environment	13
3.5.2 Common training strategy	13
3.5.3 Mask R-CNN training	14
3.5.4 YOLOv8s segmentation training	14
3.5.5 Loss functions and optimization objectives	14
3.5.6 OCR-based button interpretation	15
3.6 Evaluation metrics	16
3.6.1 YOLO segmentation evaluation	16
3.6.2 OCR evaluation	17
3.7 Embedded deployment and optimization	18
3.8 ROS pipeline architecture	18
3.8.1 Deployment platform	18
3.8.2 Node responsibilities and outputs	19

3.8.3	Distance estimation and collector design	20
4	Results	22
4.1	Dataset quality results	22
4.2	YOLO hyperparameter tuning and evaluation	23
4.2.1	Baseline YOLOv8s segmentation performance	23
4.2.2	Hyperparameter tuning results	23
4.2.3	Comparison with YOLO26	24
4.2.4	Environment-specific fine-tuning with ADC data	25
4.3	OCR hyperparameter tuning and evaluation	28
4.3.1	OCR hyperparameter tuning	28
4.3.2	Final OCR benchmark	29
4.4	Model hardware optimization results	30
4.5	ROS pipeline performance	31
4.5.1	Qualitative end-to-end pipeline output	32
5	Discussion	34
5.1	Interpretation of the main findings	34
5.2	Relation to the research questions	35
5.3	Limitations	36
5.4	Implications for future robot deployment	36
6	Conclusions and recommendations	38
6.1	Conclusions	38
6.2	Recommendations	39
6.3	Final reflection	40
A	Supplementary hyperparameter tuning curves	43
A.1	YOLO hyperparameter tuning curves	43
A.2	OCR hyperparameter tuning curves	44
B	Supplementary model evaluations	45
B.1	Mask R-CNN checkpoint progression during annotation support	45
B.2	Baseline YOLOv8s segmentation versus final Mask R-CNN	45
B.3	Reference image for the hardware optimization benchmark	46
C	Supplementary deployment experiments	47
C.1	Unfiltered single-button detector comparison	47
C.2	Single-button end-to-end confidence behavior	47
C.3	Mixed-environment replay review by distance bin	47
D	Supplementary dataset quality plots	49
D.1	Image-level quality distributions	49
D.2	Spatial button distribution	49

E	ROS pipeline parameter reference	52
E.1	YOLO inference node	52
E.2	OCR inference node	53
E.3	Active Data Collector	54
E.4	Detailed runtime timing breakdown	55
E.5	Example final pipeline output	56

1 Introduction

This chapter introduces the project context, defines the perception problem addressed in this thesis, states the project objective and research questions, and outlines the remainder of the report.

1.1 Background

Robotics in healthcare is moving from experimental pilots to practical deployment, supporting tasks such as logistics, delivery of medical supplies, and environmental assistance. These systems are intended not to replace clinical staff, but to support them by reducing repetitive physical work and enabling more time for patient care [6, 15]. As hospitals face increasing operational pressure and staff shortages, autonomous robots capable of safe navigation and interaction with the environment hold significant potential.

Hospitals are typically multi-floor environments that rely on infrastructure such as elevators, automatic doors, and access-control systems. For a robot to move freely and perform tasks across different floors, it must be able to detect and interact with these human-machine interfaces, including elevator buttons and door panels. Reliable perception of such interaction elements is therefore a key capability for achieving truly autonomous navigation in hospital settings. Without this, robots remain confined to limited single-floor operation, require costly dedicated interfaces to doors or elevators, or depend on human assistance for vertical transport.

The Ambyon ONE is a service robot currently under development by *Ambyon*, a Dutch startup focused on robotics for healthcare and logistics applications. Ambyon’s strategy centers on developing robots that do not require infrastructural changes, positioning this as a key differentiating selling point for integration into existing hospital environments.

At *HAN Automotive Research*, the CareBOTS platform is used for the integration and testing of intermediate deliverables from multiple HAN research activities within the broader CareBOTS project. Within this thesis, the platform serves as the integration and test unit for the proposed perception pipeline. While it is not intended to replicate the Ambyon ONE robot, it is equipped with a similar sensor stack and processing unit, enabling realistic simulation of Ambyon’s operational conditions.

The CareBOTS system uses an Red-Green-Blue plus Depth (RGB-D) camera and Embedded GPU platform for edge AI (NVIDIA Jetson) computing module to enable on-device



Figure 1.1: CareBOTS prototype at HAN Automotive Research

perception suitable for real-time operation in clinical environments. A key research challenge is enabling the robot to accurately detect and classify interaction points such as elevator buttons, access readers, and door panels under varying lighting, reflections, occlusions, and diverse panel designs - while maintaining efficient performance within embedded hardware constraints.

This project contributes to that research direction by designing and evaluating a scalable computer-vision workflow for elevator-button detection and interpretation on embedded systems. The work includes dataset creation, model design and optimisation, evaluation on prototype hardware, and preparation for future extensions. The resulting system will serve as a foundation for autonomous hospital mobility and support the development of assistive robotic technologies.

1.2 Problem Definition

The Amblyon ONE robot currently lacks a reliable computer-vision system to detect, interpret, and localize elevator buttons in hospital environments under embedded-hardware constraints, preventing autonomous interaction with building infrastructure.

1.3 Project Objective

Develop and validate a scalable computer-vision workflow that enables a service-robot platform to autonomously detect and classify human-interaction elements in hospital environments using embedded hardware.

1.4 Research Question(s)

Main Research Question:

How can a computer-vision workflow be designed and implemented to enable a service robot to detect and interpret human-interaction elements in hospital environments using embedded hardware?

Sub-questions:

1. What are the relevant human-interaction points in hospital environments that should be detected and classified by the computer-vision workflow?
2. What computer-vision methods and model architectures are most suitable for detecting and classifying human-interaction elements in service-robot environments?
3. How can a high-quality dataset be designed and constructed for achieving robust model performance?
4. Which model-optimization techniques support real-time inference on embedded hardware while preserving detection accuracy?
5. How can the workflow support scalable model improvements, including potential data collection and retraining during deployment?

1.5 Outline of the report

Chapter 1 introduces the project background, problem definition, objective, and research questions. Chapter 2 reviews the literature relevant to interaction-point perception, including task formulation, public datasets, annotation strategies, model design, and embedded deployment constraints. Chapter 3 explains the methodology used in this thesis, including dataset selection and preparation, the annotation workflow, model training, evaluation, deployment, and the final Robot Operating System (ROS) pipeline design. Chapter 4 presents the results of the dataset-quality assessment, model training and tuning experiments, held-out test evaluations, and deployment benchmarks. Chapter 5 interprets these results in relation to the research questions, discusses the main limitations, and considers implications for future robot deployment. Chapter 6 concludes the report and presents recommendations for future work. The appendices provide supplementary model evaluations, dataset-quality plots and additional technical details of the deployed pipeline.

2 Literature Review

This chapter reviews the scientific literature relevant to interaction-point perception in hospital environments, with a focus on instance segmentation, dataset design, annotation strategies, and deployment of computer-vision models on embedded systems.

2.1 Perception of human-machine interaction elements

Human-machine interaction elements such as elevator buttons, door panels, and access readers are important targets for indoor service robots operating in hospital environments. Previous work in healthcare and service robotics shows that reliable interaction with such infrastructure is required for autonomous operation across multiple floors and restricted areas [6, 15].

From a computer-vision perspective, these interaction elements are challenging to detect because they are small, visually similar, and often arranged closely together on control panels. Studies on visual perception for robotic interaction report that elevator panels typically contain many buttons with minimal visual differences, which makes robust localization difficult under varying lighting conditions, reflections, and viewing angles [2].

General indoor-robot perception systems often focus on detecting large objects or structural elements and assume that interaction with building infrastructure is externally supported or environment-specific. However, research on autonomous service robots emphasizes that scalable deployment requires perception systems that can directly identify actionable interface elements from camera and sensor data, without relying on prior knowledge of the environment [24].

These findings indicate that interaction-point perception is a fine-grained vision problem in which individual button instances on a panel must be localized accurately within visually dense scenes. At the same time, the cited studies mainly motivate the problem rather than solving it under the embedded and deployment-oriented constraints of this project. This observation motivates the need for vision methods that go beyond coarse detection of panels or regions and instead support precise identification of individual button instances.

2.2 Formulating interaction-point perception as a vision problem

Vision-based perception problems in robotics are commonly formulated as image classification, object detection, or semantic segmentation tasks. Each formulation implies different assumptions about the structure of the scene and the required output, which directly affects suitability for robotic interaction.

Image classification assigns a single label to an entire image and is therefore unsuitable for scenarios where multiple interaction elements appear simultaneously. Object detection extends this formulation by predicting bounding boxes and class labels for individual objects; in this project, those objects correspond to individual button instances on an elevator panel. While detection-based approaches have been widely used for indoor perception tasks, bounding boxes

provide only coarse spatial localization and do not capture the precise shape of interaction elements [19, 18].

Semantic segmentation addresses this limitation by assigning a class label to each pixel in the image. Formally, semantic segmentation can be expressed as a pixel-wise classification function

$$f_{\theta} : \mathbb{R}^{H \times W \times C} \rightarrow \{1, \dots, K\}^{H \times W},$$

where H and W denote image height and width, C the number of input channels, and K the number of semantic classes [13]. Although this formulation enables precise localization, all pixels belonging to the same class are treated as a single region. As a result, semantic segmentation cannot distinguish between multiple instances of the same object class.

For interaction-point perception, this limitation is critical. Elevator panels often contain many visually similar button instances, each corresponding to a distinct action [2]. By definition, semantic segmentation assigns a class label to each pixel but does not separate same-class instances [13], whereas instance segmentation predicts a separate mask for each detected instance [5]. For an elevator panel, a semantic segmentation output would therefore label button pixels by class without separating adjacent buttons into distinct actionable targets. This lack of instance-level differentiation makes semantic segmentation unsuitable for downstream decision-making and physical interaction, where each button must be treated as a separate actionable element.

Instance segmentation extends semantic segmentation by jointly performing object detection and pixel-wise classification, producing a separate mask for each object instance. This formulation can be expressed as a set of instance masks

$$\mathcal{M} = \{M_1, M_2, \dots, M_N\},$$

where each $M_i \in \{0, 1\}^{H \times W}$ represents a binary mask for one object instance [5]. Instance segmentation therefore enables both precise spatial localization and differentiation between multiple objects of the same class. Figure 2.1 illustrates the conceptual differences between image classifi-




Figure 2.1: Comparison of image classification, object detection, semantic segmentation, and instance segmentation. Adapted from Liu [11].

cation, object detection, semantic segmentation, and instance segmentation for visual perception tasks.

Based on findings in the literature and the requirements of interaction-point detection, instance segmentation provides the most appropriate problem formulation for this project. It allows individual buttons to be identified as distinct entities while preserving pixel-level accuracy, which is essential for reliable robotic interaction in dense and visually complex control interfaces. The literature therefore clarifies the correct task formulation, but not yet which model family best satisfies this requirement on embedded robotic hardware.

2.3 Public datasets and annotation strategies.

Public datasets for vision-based detection of elevator interfaces are limited, as most indoor robotics datasets focus on navigation or scene understanding rather than interaction elements [21]. This limits their direct applicability to interaction-point perception tasks.

The IRSO 2018 dataset provides Red–Green–Blue (RGB) images of elevator panels captured in realistic indoor environments and was originally introduced for object detection tasks [14]. Due to its focus on real elevator interfaces and visual diversity, it is well suited for studying interaction-point perception in service robotics.

Recent work reports a large-scale dataset with pixel-wise annotations for elevator button segmentation [12]. However, despite being described in the literature, the corresponding semantic segmentation dataset is not publicly available at the time this project is made and therefore cannot be used.

To enable instance-level interaction-point detection, this work reuses images from the IRSO 2018 dataset and replaces the original annotations, namely Extensible Markup Language (XML)-formatted bounding boxes and textual labels, with instance segmentation masks. The new annotations are stored in the Common Objects in Context (COCO) format, which supports instance-level labeling and is widely used across computer vision frameworks, facilitating reuse and future extensions [10]. This reveals a practical gap in the literature: relevant public images exist, but publicly reusable instance-segmentation annotations for elevator interfaces are still limited.

2.4 Model design principles, evaluation metrics, and semantic interpretation

Instance segmentation models for robotic perception typically follow a modular design in which feature extraction, object localization, and pixel-level mask prediction are handled by separate components within a single network. This design enables shared visual representations while supporting accurate instance-level localization, which is required for interaction-point perception in dense control interfaces.

Region-based instance segmentation models such as Mask Region-Based Convolutional Neural Network (Mask R-CNN) achieve high localization accuracy through a two-stage design in which candidate regions are proposed first and then refined by classification and mask-prediction heads [5]. This architecture is widely used as a reference and benchmark in instance-

segmentation research because of its strong instance-level accuracy. However, its multi-stage processing increases computational cost and memory use, which limits suitability for real-time deployment on embedded robotic platforms.

To address these constraints, single-stage instance-segmentation architectures have been developed that combine detection and mask prediction in a unified pipeline. You Only Look Once (YOLO)-based segmentation models, including the YOLOv8s segmentation model, remain CNN-based vision models, but differ from Mask R-CNN in that they avoid a separate region-proposal stage and are optimized for faster end-to-end inference [8]. In practice, this design trades some flexibility and, in some cases, peak segmentation accuracy for substantially improved inference speed and computational efficiency, which makes such models attractive for embedded robotic deployment.

Model performance is commonly evaluated using overlap-based metrics that measure agreement between predicted masks and ground-truth annotations. Intersection over Union (IoU) is defined as

$$\text{IoU} = \frac{|M_p \cap M_{gt}|}{|M_p \cup M_{gt}|},$$

where M_p and M_{gt} denote the predicted and ground-truth masks. Average Precision (Average Precision (AP)), computed over multiple IoU thresholds, is widely used as a standard evaluation metric for instance segmentation tasks [10].

For elevator-button perception, instance localization can be followed by a recognition stage that determines the functional meaning of each detected button. This can be achieved using character recognition, Optical Character Recognition (OCR), or dedicated symbol classification models applied to isolated button regions [12]. Accurate instance segmentation is a critical enabling step for such approaches, as it provides clean, well-aligned inputs for subsequent classification or text recognition.

In this work, instance segmentation is treated as the primary perception task, as reliable instance-level localization is a prerequisite for any semantic interpretation of interaction elements. The resulting pipeline is therefore designed to support integration of OCR or custom button classifiers, enabling the robot to associate detected buttons with their intended functions in downstream processing stages. The literature gives clear architectural candidates, but the trade-off between multi-stage accuracy and single-stage deployability remains an open design decision for this specific use case.

2.5 Embedded deployment constraints in robotic perception

Vision-based perception systems deployed on mobile service robots must operate under strict computational and energy constraints. Prior work in robotic vision reports that deep-learning models deployed on embedded platforms are limited by onboard memory, processing power, and power consumption, which directly affects real-time performance [3].

Studies evaluating deep-learning inference on NVIDIA Jetson-class hardware show that models achieving high accuracy on desktop Graphics Processing Units (GPUs) often exhibit increased latency and memory usage when deployed on embedded devices [9]. These findings indicate that benchmark accuracy alone is not a sufficient criterion for model selection in robotic perception

systems.

For instance segmentation tasks, this trade-off is particularly relevant, as mask prediction introduces additional computational overhead compared to bounding-box detection. Research on efficient model design highlights that architectural choices such as backbone depth and feature reuse have a significant impact on inference efficiency [22]. Based on these insights, this project treats model efficiency and deployability as primary design constraints alongside segmentation accuracy.

Consequently, instance segmentation models considered in this work are evaluated not only with respect to localization performance, but also in terms of their suitability for embedded deployment. However, relatively few studies report this trade-off for elevator-interface perception specifically. This literature-driven perspective directly informs the model selection and optimization strategies described in Chapter 3 (Methodology).

2.6 Scalability, robustness, and reproducibility.

Several studies emphasize that robust perception systems require more than a single well-performing model. Instead, scalability is achieved through structured data management, consistent annotation practices, and reproducible training pipelines that support incremental dataset expansion and model updates [4]. These practices reduce sensitivity to domain-specific bias and enable adaptation to new environments without redesigning the entire system.

From these findings, this project adopts a pipeline-oriented approach in which dataset creation, annotation, training, and evaluation are treated as modular and repeatable processes. Rather than optimizing a single model for a fixed dataset, the focus is placed on maintaining a perception workflow that can be extended with new data and retrained as deployment conditions evolve. Existing literature strongly advocates reproducibility, but gives less attention to practical feedback loops for environment-specific data collection and retraining after deployment.

2.7 Summary and identified research gap

The reviewed literature shows that reliable perception of interaction elements in robotic systems requires instance-level localization, suitable annotation strategies, and awareness of embedded deployment constraints. Existing studies demonstrate the technical feasibility of these components, but often address them in isolation.

At the same time, publicly available datasets for elevator interfaces are limited, and existing annotations are commonly designed for object detection rather than segmentation tasks. In addition, many approaches focus on model performance without explicitly considering scalability, reproducibility, or deployment on embedded robotic platforms.

This project addresses these gaps by repurposing a realistic elevator dataset with instance-level annotations and by developing a scalable and deployment-aware computer-vision pipeline for elevator-button perception in hospital environments. The following chapter describes the methodology used to realize this approach.

3 Methodology

This chapter describes the methodology used to develop a computer-vision pipeline for detecting and interpreting elevator buttons in hospital environments. The project is conducted with the objective of building a scalable and deployment-oriented perception system that can be iteratively improved across different hospital settings and reused within a robot fleet.

3.1 Overview of the proposed perception workflow

The proposed methodology focuses on instance-level localization of interaction elements such as elevator buttons and control panels. Instance segmentation is used to detect individual button instances at pixel level, allowing precise localization and separation of closely spaced elements. This capability provides the basis for subsequent interpretation of button functions and autonomous interaction.

The methodology is designed to support an iterative training and deployment workflow. An initial base model is deployed on the robot, where it performs interaction-point perception and collects visual data during operation. Selected data can be transferred to an offboard system, where annotations may be reviewed, corrected, or extended before retraining the perception models. The updated models are then redeployed to the robot and become the new default for continued operation.

In the implemented improvement loop, not every observed frame is exported for later review. Instead, the Active Data Collector (ADC) applies quality gates and stores only frames that contain accepted button targets satisfying detector and OCR confidence thresholds, and compliance with the additional export constraints described later in this chapter, such as image-quality, stability, and de-duplication checks. The offboard review stage therefore starts from a filtered set of high-value samples rather than from the full camera stream, which makes the improvement loop in Figure 3.1 an implemented data-collection process rather than only a conceptual workflow diagram.

This workflow enables gradual adaptation to new hospital environments and allows knowledge acquired in one deployment to be reused and refined in subsequent deployments. The target end-to-end training and deployment process developed in this project is illustrated in Figure 3.1. Figure 3.1 therefore serves as a combined workflow-and-architecture overview: it links the source dataset, annotation and training stages, deployment on the robot, and the feedback path introduced by the ADC. The training-side annotation architecture is detailed later in Figure 3.3, and the deployed runtime architecture is detailed in Figure 3.6.

The following sections describe each stage of the methodology in detail, beginning with dataset selection and preparation.

3.2 Dataset selection and preparation

Within the overall workflow-and-architecture overview shown in Figure 3.1, this section covers the preparation steps that precede the first model-training phase. Before any segmentation or




Figure 3.1: Overall workflow-and-architecture overview of the proposed perception system.

recognition model can be trained, the source dataset must be selected, split, reformatted, and extended where needed for the later deployment setting.

This project uses the elevator-panel dataset released by Liu et al. [12]. For this thesis, the working dataset comprised approximately 2000 images, split into 1122 training images, 281 validation images, and 602 held-out test images. The test split was reserved for final evaluation only.

The original dataset provided XML annotations with bounding boxes and textual labels, but no instance masks. To support instance segmentation, the images were re-annotated in a semi-automatic Computer Vision Annotation Tool (CVAT) workflow using a Mask R-CNN pre-annotation model, as described in Section 3.3. The resulting annotations were stored in COCO-style JavaScript Object Notation (JSON) and then converted to the label format required by YOLO-based segmentation training.

In the later deployment-oriented adaptation stage, the dataset was extended with 46 additional images collected in Han Building R29, the pilot environment for robot deployment testing. These images were added only after the baseline and hyperparameter tuning stages, so that the initial YOLO experiments remained directly comparable with the original benchmark setting.

The resulting dataset structure provides a consistent and reusable foundation for annotation, model training, and evaluation, as described in the following sections.

In addition to the instance-segmentation datasets, a recognition dataset for the semantic interpretation stage was derived from the same annotated source images. For this purpose, each annotated button instance was cropped from the original image using its existing bounding-box annotation, with an additional contextual padding of 15% of the box width and height. This produced a button-centered crop dataset for OCR fine-tuning without requiring any additional



Figure 3.2: Example raw image from the dataset.

manual image collection.

The OCR dataset builder grouped crops by source image before creating a fixed 80/20 train/validation split with a deterministic random seed. As a result, all button crops originating from one panel image were placed either in the training set or in the validation set, but never in both. This prevents information leakage caused by highly similar button appearances within the same elevator panel. The held-out OCR test set was generated directly from the original test images and their annotations.

Several OCR dataset variants were explored during development. The final variant used for the main recognition experiments removed labels named exactly `text` and labels starting with the prefix `text_`. These labels do not encode the actual button content, but merely indicate the presence of unspecified free text. Retaining them caused the recognizer to over-predict generic “TEXT”-like outputs and reduced generalization to unseen textual button content. The final OCR dataset, contained 11,953 training crops, 2,973 validation crops, and 6,861 held-out test crops.

3.3 Annotation strategy and workflow

Instance-level ground truth was created using a semi-automatic workflow that combines model-based pre-annotation with manual verification. This strategy was used to efficiently generate high-quality instance segmentation annotations while maintaining consistency across the dataset.

Annotations were created in CVAT [20], deployed locally using Docker [16] within a Windows Subsystem for Linux 2 (WSL2)-based Linux environment on a Windows host. Automatic annotation was enabled through CVAT–Nuclio integration, where the model was deployed as a Nuclio serverless function [7]. The instance segmentation model used for automatic annotation was Mask R-CNN [5].

During automatic annotation, CVAT sends task images to the Nuclio function, where Mask R-CNN predicts binary masks for each detected button instance. These masks are converted into polygon shapes within the Nuclio function and returned to CVAT as annotation data. The overall workflow used in this project is shown in Figure 3.3.

To support iterative improvement, the dataset was annotated in batches of approximately 150 images. After each three batches, Mask R-CNN annotation-support model was retrained using the newly annotated data and reused for subsequent annotation rounds. Model check-




Figure 3.3: Semi-automatic annotation workflow on a Windows host using WSL2, Docker, CVAT, and Nuclio.

points were created after approximately 90, 449, and 900 annotated images, followed by a final checkpoint trained on the full 1400-image training split. This enabled progressively improved pre-annotations and reduced manual correction effort. A quantitative comparison of these annotation-support checkpoints on the held-out test set is provided in Appendix B.

All automatically generated annotations were also manually reviewed and corrected when needed to ensure accurate instance boundaries, correct separation of adjacent buttons, and consistent labeling across both the training and test sets. This manual verification step was retained because fully automatic acceptance would have been faster, but it would also have propagated model mistakes into the training data, especially for touching buttons, incomplete or missed masks. A fully manual workflow, by contrast, would have offered maximum control at substantially much higher annotation time and lower scalability. The chosen hybrid approach therefore balanced annotation efficiency with dataset quality.



Figure 3.4: Example of an annotated image showing instance segmentation masks.

3.4 Dataset quality assessment

A brief dataset quality assessment was performed before the main training experiments to check whether the dataset contained a substantial subset of clearly unusable images.

For each training image, two basic image-level indicators were computed:

- a blur indicator based on the variance of the Laplacian, used only as a relative sharpness measure.
- the mean grayscale intensity, used as a simple brightness measure.

In addition, the existing instance annotations were used to derive basic geometry statistics, including the approximate button size in pixels, the relative image area occupied by button instances, and the normalized button location within the image. These geometry measurements were used to check whether the training set covered a realistic range of viewing distances and camera framing conditions.

The quantitative indicators were not used as automatic pruning rules. Instead, images with low sharpness values, very dark appearance, or unusual geometric properties were manually reviewed. In addition, a model-assisted assessment was run over the full training split using a baseline instance-segmentation model. For every training image, the model produced per-image mean IoU, false-positive count, and false-negative count. These metrics were used to identify the lowest-IoU and highest-error cases for manual inspection. The model-assisted pass therefore covered the complete training set, while the manual review focused on the most difficult cases highlighted by these per-image metrics. The descriptive outcomes of this assessment are reported in Section 4.1.

3.5 Model architectures and training strategy

Three model families are employed in this project, each serving a distinct role within the overall perception pipeline. Mask R-CNN is used as an annotation-support model in the semi-automatic labeling workflow, YOLOv8s segmentation is trained and evaluated as the deployment-oriented instance segmentation model, and the OCR recognizer is used for semantic interpretation of the detected button crops.

3.5.1 Experimental hardware and software environment

All models training experiments reported in this thesis were executed on the same workstation. The host system ran Ubuntu 22.04.5 LTS with Linux kernel 6.8.0-90. The machine was equipped with an AMD Ryzen 9 9950X 16-core processor and an NVIDIA GeForce RTX 4090 graphics card with 24 GiB of Video Random Access Memory (VRAM). The training software environment used Python 3.10.12 and PyTorch 2.5.1 compiled against Compute Unified Device Architecture (CUDA) 12.1. This configuration provided the computational platform for the baseline training, hyperparameter tuning and the comparative evaluation runs discussed in Chapters 3 and 4.

3.5.2 Common training strategy

All trainable vision models were fine-tuned from pretrained weights rather than trained from scratch, and their task-specific output heads were adapted to the project datasets. For transfer learning, YOLOv8-Seg was initialized from the official Ultralytics `yolov8s-seg.pt` checkpoint pretrained on COCO for instance segmentation, Mask R-CNN from Torchvision’s default COCO-pretrained `MaskRCNN_ResNet50_FPN` weights, and the OCR recognizer from the official PaddleOCR `PP-OCRv5_mobile_rec` pretrained recognition checkpoint. Across the experiments, the train/validation split was used for model fitting and model selection, while the held-out test split was reserved for final comparison only.

3.5.3 Mask R-CNN training

Mask R-CNN was used only as an annotation-support model within the semi-automatic CVAT workflow. It was trained on COCO-formatted instance annotations and retrained iteratively as the annotated dataset grew. The checkpoint progression and held-out reference comparison are summarized in Appendix B, since only the role of Mask R-CNN in the annotation workflow is relevant to the main methodology.

3.5.4 YOLOv8s segmentation training

The common training strategy was applied to the deployment-oriented YOLOv8s segmentation model using the Ultralytics framework [23]. The initial baseline configuration was scheduled for up to 100 epochs with an input resolution of 640×640 , dynamic batch sizing (`batch=-1`), warmup over 5 epochs, and early stopping patience of 10 epochs. The augmentation policy followed the standard Ultralytics recipe used throughout this study, including small geometric transformations, horizontal flipping, and mosaic augmentation.

Although this baseline produced strong results, its training settings were selected empirically. A structured hyperparameter tuning procedure was therefore conducted to determine whether performance could be improved in a more systematic and reproducible way. All tuning runs were logged and tracked in Weights & Biases (W&B), which was used to store the hyperparameter configurations, training curves, and validation metrics of individual experiments.

The baseline settings were chosen as a pragmatic starting point. The 640×640 input size balanced preservation of small-button detail against the training and inference costs expected for embedded deployment, although lower resolution can make small or distant button instances harder to localize accurately. Dynamic batch sizing, warmup, and early stopping were used to stabilize optimization and use the available GPU memory efficiently.

The three stages of the hyperparameter tuning procedure are summarized in Table 3.1. The table is intended to provide a compact overview of the explored search spaces and stopping criteria.

After this tuning process, two further experiments were performed: a comparison with YOLO26-small and a domain-adaptation fine-tuning run using the mixed dataset that included images collected in the target deployment environment.

3.5.5 Loss functions and optimization objectives

Both instance-segmentation models were optimized with composite losses that jointly account for classification, localization, and mask prediction.

Mask R-CNN. The training objective of Mask R-CNN is defined as a multi-task loss composed of three terms [5]:

$$\mathcal{L}_{\text{Mask R-CNN}} = \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{box}} + \mathcal{L}_{\text{mask}}, \quad (3.1)$$

where $\mathcal{L}_{\text{cls}}$ denotes the classification loss, implemented as categorical cross-entropy; $\mathcal{L}_{\text{box}}$ represents bounding-box regression using the smooth L_1 loss; and $\mathcal{L}_{\text{mask}}$ is a pixel-wise binary cross-entropy loss applied to each predicted instance mask.

YOLOv8s segmentation. YOLOv8s segmentation integrates detection and segmentation into

Table 3.1: Three-stage hyperparameter tuning procedure for YOLOv8s segmentation.

Stage	Purpose	Tuned or varied settings	Training setup	Selection rule
Stage 1	Narrow hyperparameter search around the baseline recipe	<code>lr0</code> , <code>lrf</code> , <code>weight_decay</code> , <code>momentum</code> , <code>hsv_h</code> , <code>hsv_s</code> , <code>hsv_v</code> , <code>degrees</code> , <code>translate</code> , <code>scale</code> , <code>mosaic</code> , <code>fliplr</code> , <code>warmup_epochs</code> , <code>close_mosaic</code>	24 trials, 30 epochs each; sanity run of 12 epochs before-hand	Top runs by validation Mean Average Precision (mAP)
Stage 2	Capacity sweep using the best Stage 1 configurations	Batch size $\in \{8, 16\}$, image size $\in \{640, 768\}$; Stage 1 parameters kept fixed	100 epochs, patience 15	Top 5 Stage 1 runs with $\text{mAP}_{50:95}^{(M)} \geq 0.30$
Stage 3	Late-stage refinement from the best Stage 2 checkpoints	Resume from best checkpoint, halve <code>lr0</code>	180 epochs	Top 2 Stage 2 runs with $\text{mAP}_{50:95}^{(M)} \geq 0.50$

a single-stage architecture. The total training loss is expressed as a weighted sum of classification, localization, and segmentation losses [8]:

$$\mathcal{L}_{\text{YOLO}} = \lambda_{\text{cls}}\mathcal{L}_{\text{cls}} + \lambda_{\text{box}}\mathcal{L}_{\text{box}} + \lambda_{\text{seg}}\mathcal{L}_{\text{seg}}, \quad (3.2)$$

where $\mathcal{L}_{\text{seg}}$ represents the segmentation loss applied to the predicted mask outputs, and the weighting coefficients λ control the relative contribution of each task during optimization.

3.5.6 OCR-based button interpretation

Following instance segmentation, detected button regions are passed to a semantic interpretation stage. Segmented button instances are cropped and processed by an OCR-based recognition module to extract character labels or symbols associated with each button. Because button localization is already solved upstream by the segmentation model, the recognizer does not need to detect text in the full image and instead operates on tightly bounded button crops. This reduces computational cost and enables the robot to associate localized interaction points with their functional meaning during downstream decision-making.

Recognition model architecture. For cropped-button recognition, a lightweight OCR recognizer based on Paddle Optical Character Recognition (PaddleOCR) [17] was selected. This model was preferred because it is designed for efficient deployment and already provides a compact recognition backbone suitable for embedded hardware. Only the recognition component of the OCR stack was used. The detection component of PaddleOCR was intentionally omitted from the final recognition pipeline, because button localization was already solved upstream by the instance segmentation model. This avoids redundant computation and makes the OCR stage conceptually cleaner: segmentation localizes the button instance, and OCR interprets that localized instance.

Fine-tuning strategy. The recognizer was initialized from the official pretrained English

PP-OCRv5_mobile_rec weights and fine-tuned on the project-specific crop dataset described in Section 3.2. During training, each crop was resized to the fixed input resolution of 48×320 defined in the PaddleOCR recognition configuration so that all samples reached the network with a consistent image height and width. Optimization used Adam with cosine learning-rate scheduling, warmup, and L2 regularization. Recognition-specific augmentation and multiscale sampling were enabled to improve robustness and reduce overfitting to a single visual presentation of the button crops.

The implementation followed the standard PaddleOCR training workflow [17], using the official `tools/train.py` and `tools/eval.py` scripts with project-specific YAML Ain’t Markup Language (YAML) configuration files.

The OCR tuning procedure is summarized in Table 3.2. Stage 1 was used to determine an appropriate epoch budget, while Stage 2 fixed that budget and performed a narrow optimizer sweep. All OCR tuning runs were logged in W&B so that training curves, hyperparameter choices, and validation metrics could be compared systematically.

Table 3.2: Overview of the staged OCR hyperparameter tuning procedure.

Stage	Purpose	Parameters varied	Training budget	Selection criterion
Stage 1	Epoch selection under fixed baseline optimizer settings	Epochs $\in \{30, 50, 60\}$	One run per epoch budget	Best evaluation accuracy, supported by normalized edit distance
Stage 2	Narrow optimizer sweep at the selected epoch budget	Learning rate, weight decay, and warmup	Fixed at 50 epochs	Best evaluation accuracy, supported by normalized edit distance

3.6 Evaluation metrics

The evaluation in this project follows the metric outputs of the training and validation toolchains used in the experiments. YOLO-based segmentation is reported with COCO-style box and mask metrics, while the OCR recognizer is reported with exact-match accuracy and normalized edit distance. Validation dataset was used for model selection and tuning, and held-out test dataset was reserved for final reporting. The Mask R-CNN annotation-support comparison is included only as supplementary material in Appendix B.

3.6.1 YOLO segmentation evaluation

The YOLO experiments report the standard Ultralytics/COCO validation and test metrics for both bounding boxes and segmentation masks. The main metric used for model selection is mask mAP@[0.50:0.95] . This metric averages AP over IoU thresholds from 0.50 to 0.95 in increments of 0.05 and is the most representative single indicator of segmentation quality. In addition, `metrics/mAP50(M)` is reported to show performance at the less strict 0.50 IoU threshold. This reporting convention follows the standard COCO-style Ultralytics metrics interface [10, 23].

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN} \quad (3.3)$$

$$\text{mAP}_{50:95} = \frac{1}{10} \sum_{t \in \{0.50, 0.55, \dots, 0.95\}} \text{AP}_t \quad (3.4)$$

The complete set of reported YOLO metrics in this report therefore includes mask precision, mask recall, $\text{mAP}_{50}(\text{M})$, and $\text{mAP}_{50-95}(\text{M})$, together with the corresponding box metrics $\text{precision}(\text{B})$, $\text{recall}(\text{B})$, $\text{mAP}_{50}(\text{B})$, and $\text{mAP}_{50-95}(\text{B})$. These additional values are retained because they help explain the behavior of a model beyond a single summary number. For example, precision and recall indicate whether a model tends to miss valid instances or to produce false positives, while the box metrics make it possible to distinguish localization quality from mask quality. However, when ranking segmentation models in this report, mask $\text{mAP}@[0.50:0.95]$ remains the primary comparison criterion.

3.6.2 OCR evaluation

The OCR recognition models are evaluated using exact-match accuracy and normalized edit distance. Exact-match accuracy is the fraction of button crops for which the predicted label string matches the ground-truth label exactly after the same normalization procedure is applied to both. This metric is strict and is therefore suitable for the final semantic interpretation task, where predicting 4 instead of 14, or OPE instead of OPEN, must be counted as an error.

$$\text{Accuracy}_{\text{OCR}} = \frac{N_{\text{correct}}}{N_{\text{total}}} \quad (3.5)$$

where N_{correct} is the number of exactly correct predictions and N_{total} is the total number of evaluated crops.

In addition to exact-match accuracy, normalized edit distance is used to quantify how close an incorrect prediction is to the target string. This metric is derived from the Levenshtein edit distance between prediction and target and is normalized to the interval $[0, 1]$, where 1 indicates a perfect match. Unlike exact-match accuracy, normalized edit distance gives partial credit to near-correct predictions, which is useful when analyzing OCR behavior on short strings and icon labels that are often confused by only one or two characters.

$$\text{NED} = 1 - \frac{1}{N} \sum_{i=1}^N \frac{d_{\text{edit}}(\hat{y}_i, y_i)}{\max(|\hat{y}_i|, |y_i|, 1)} \quad (3.6)$$

where $\hat{y}_i$ is the predicted string, y_i is the corresponding ground-truth string, and d_{edit} denotes the Levenshtein edit distance.

During OCR tuning, validation accuracy was treated as the primary model-selection metric, while normalized edit distance was used as a secondary quality indicator. Final OCR benchmarking was then performed on the held-out OCR test split using the official PaddleOCR evaluation script with the same configuration family as used during training [17].

3.7 Embedded deployment and optimization

To enable real-time execution on embedded hardware, the trained models are optimized and deployed on a NVIDIA Jetson Orin platform. This section corresponds to the deployment branch of the overall workflow-and-architecture overview introduced in Figure 3.1. The deployment workflow is designed to be identical for both the instance segmentation model and the OCR-based button classification model.

Models are first trained in PyTorch, exported to the Open Neural Network Exchange (Open Neural Network Exchange (ONNX)) [1] format, and then converted on the NVIDIA Jetson Orin device into optimized TensorRT (TensorRT) engines. Engine generation was performed directly on the target hardware to ensure compatibility with the device-specific GPU architecture. Half-precision floating point format (FP16) precision was used to reduce memory footprint and improve inference speed.

The final TensorRT `.engine` models were validated directly on the target device by running them on a subset of held-out test images and checking that the predictions were numerically sensible and visually consistent with the trained desktop models.

The resulting TensorRT engines are integrated into ROS nodes responsible for perception tasks. At runtime, image data from the RGB-D camera is passed through the ROS pipeline and processed by the TensorRT engines to produce instance segmentation masks and corresponding button labels.

The complete deployment and optimization workflow is illustrated in Figure 3.5. This approach ensures a consistent and efficient deployment pipeline and enables fair evaluation of different model architectures under identical runtime constraints.




Figure 3.5: Jetson deployment workflow.

3.8 ROS pipeline architecture

After deployment-oriented model selection, the TensorRT engines were integrated into a modular ROS pipeline on the target robot computer. This stage links the trained perception models to live camera input, downstream message passing, and deployment data collection.

3.8.1 Deployment platform

The runtime pipeline was executed on an NVIDIA Jetson Orin Nano running Ubuntu 22.04.5 LTS and ROS Humble. The device used JetPack 6.2.1, corresponding to NVIDIA Linux for Tegra release R36.4.7 with Linux kernel 5.15.148-tegra. TensorRT 10.7.0 was used for optimized

inference. Visual input was provided by a Stereolabs ZED X RGB-D camera connected through a ZED Link Duo capture card, using ZED SDK 5.2.1.

3.8.2 Node responsibilities and outputs

The deployed pipeline consists of the ZED camera driver, a YOLO inference node, an OCR inference node, and an ADC. As illustrated in Figure 3.6, the camera provides the colour image, registered depth image, and calibration data required for geometric perception. The `yolo_engine_node` converts these inputs into button candidates that already include image-plane localization through a mask and bounding box, a representative center pixel, and a depth-based camera-relative target estimate. The `ocr_trt_node` then adds semantic information by assigning the recognized button text and its confidence to each accepted candidate. Finally, the ADC receives the enriched button candidates together with the camera image and stores only those samples that pass the deployment quality gates.

The split-node design was chosen deliberately. Earlier experiments showed that placing YOLO and OCR TensorRT inference in a single Python process increased the risk of CUDA-context conflicts and made the system harder to debug. Separating the geometric perception stage from the semantic recognition stage improved modularity and allowed both runtime components to be benchmarked independently. In Figure 3.6, the gray boxes denote ROS topics, whereas the white boxes denote processing nodes or terminal outputs. The abbreviations YDBG and ODBG refer to the debug-image outputs of the YOLO and OCR nodes, respectively, while OUT denotes the final OCR-enriched button-target output stream.

Table 3.3: Main nodes in the deployed ROS perception pipeline and the information exchanged between them.

Node	Outputs used downstream	Role in the pipeline
<code>zed_wrapper</code>	Colour images, registered depth images, and camera calibration data	Provides the synchronized sensor streams used by the perception pipeline.
<code>yolo_engine_node</code>	Button candidates with mask- and bounding-box-based image localization, representative center pixels, detection scores, and depth-based camera-relative target estimates	Runs TensorRT instance segmentation, filters detections, and performs depth-based spatial projection.
<code>ocr_trt_node</code>	OCR-enriched button candidates with recognized text and text confidence, in addition to the geometric fields from the YOLO stage	Crops button regions from the current image and adds semantic label information to the accepted button candidates.
<code>data_collector</code>	Curated deployment samples containing images, annotations, and session metadata	Stores accepted field samples for later dataset expansion and model retraining.




Figure 3.6: Deployed ROS runtime architecture with the ZED X RGB-D camera, YOLO node, OCR node, and active data collector. Gray boxes indicate ROS topics, and white boxes indicate nodes or terminal outputs.

3.8.3 Distance estimation and collector design

The internal interface between the perception nodes is a custom ROS button-target message, defined for this project to bundle the geometric output of the YOLO stage and, after the OCR step, the recognized button text and text score. Image topics use a best-effort, keep-last quality-of-service policy so that the pipeline prioritizes the newest available camera frames. In practice, occasional frame drops are acceptable because processing delayed images would be less useful than continuing with recent sensor data.

Within the `yolo_engine_node`, target depth is estimated only from valid depth pixels that fall inside the predicted segmentation mask. Let $\mathcal{V}$ denote the set of valid masked depth pixels after finite-value filtering, depth range filtering, and the minimum-valid-pixel test. The representative target depth is then computed as

$$z = \text{Percentile}(\{d(i, j) \mid (i, j) \in \mathcal{V}\}, p), \quad (3.7)$$

where p is the configured depth percentile. In the current implementation, $p = 50$ is used, corresponding to the median of the valid masked depth samples. Using a percentile-based estimate rather than a plain mean reduces sensitivity to outliers caused by missing returns and reflective elevator surfaces.

The representative image location is also derived from the valid masked depth pixels. Here, i denotes the image row index and j denotes the image column index in the depth image. In

the current implementation, the pixel coordinates are taken as the median valid depth location:

$$u = \text{median}(\{j \mid (i, j) \in \mathcal{V}\}), \quad v = \text{median}(\{i \mid (i, j) \in \mathcal{V}\}). \quad (3.8)$$

If the camera intrinsics are given by focal lengths (f_x, f_y) and principal point (c_x, c_y) , the 3D target position in the camera optical frame is computed as

$$x = \frac{(u - c_x)z}{f_x}, \quad y = \frac{(v - c_y)z}{f_y}. \quad (3.9)$$

This projection step provides the button target that can be used by downstream robotic modules and by the ADC.

To support iterative improvement after deployment, the ADC stores selected images and annotations in a COCO-compatible structure. Sample export is not performed for every frame. Instead, frames pass a sequence of quality gates before they are written to a dataset session. The threshold values for image brightness, image sharpness, and minimum visible button size were derived from the dataset quality assessment in Section 4.1, where the observed brightness, blur, and button-size distributions were analyzed to define practical acceptance ranges for deployment collection. At image level, the collector therefore applies brightness and sharpness checks, while at target level it applies minimum score and size thresholds, followed by temporal stability checks and duplicate rejection.

4 Results

This chapter presents the quantitative and qualitative results obtained for the proposed elevator-interaction pipeline. It first summarizes the dataset quality assessment, then reports the results of the YOLO-based button-instance segmentation system and the OCR-based button interpretation module, followed by hardware-optimization benchmarks and integrated ROS pipeline results.

4.1 Dataset quality results

Before the main model comparison experiments, a brief dataset quality assessment was carried out on the training split to determine whether obvious image-quality problems justified dataset pruning. In total, 1403 training images containing 13,309 annotated button instances were assessed. Table 4.1 summarizes the main descriptive results.

Table 4.1: Compact summary of the dataset quality assessment on the training split.

Indicator	Summary values	Interpretation
Blur indicator (variance of Laplacian)	10th percentile = 16.53, median = 77.91, 90th percentile = 318.13	This is a relative sharpness indicator: values around 10–20 correspond to visibly soft images, while values above roughly 100 correspond to clearly sharper images. The spread therefore indicates both blurry and sharp samples.
Mean grayscale intensity	10th percentile = 55.50, median = 107.11, 90th percentile = 138.46	This follows the standard 0–255 grayscale range: values below about 50 are dark, around 100 are mid-brightness, and values above about 150 are relatively bright. The dataset therefore spans multiple illumination conditions.
Button size (equivalent diameter)	1st percentile = 19.32 px, 5th percentile = 26.12 px, median = 55.98 px	Values near 20 px represent very small targets that are difficult to detect reliably, whereas values around 50–60 px correspond to more comfortably visible buttons. Most instances are therefore larger than the hardest lower-tail cases.
Button centroid region	$x \in [0.297, 0.711]$, $y \in [0.162, 0.886]$ for the 5–95% range	The coordinates are normalized to the range [0,1], where 0.5 is near the image center and values near 0 or 1 lie close to the borders. Most buttons therefore appear in a central field-of-view band rather than at the extreme edges.

The assessment indicates realistic variation in focus, lighting, scale, and framing, but not a dominant subset of clearly unusable images. Manual review of flagged samples mainly identified

difficult yet valid cases rather than corrupt images or systematic annotation errors. No general quality-based pruning rule was therefore introduced, and the full training set was retained. Supplementary plots from this assessment are provided in Appendix D.

4.2 YOLO hyperparameter tuning and evaluation

This section presents the experimental results for the YOLO-based button-instance segmentation stage. It first reports the baseline model performance, then summarizes the staged hyperparameter tuning experiments, and finally evaluates the effect of environment specific fine-tuning with deployment collected ADC data.

4.2.1 Baseline YOLOv8s segmentation performance

The baseline YOLOv8s segmentation model established the main reference point for all subsequent experiments. On the validation split, the best checkpoint was obtained at epoch 49. This model achieved a mask mAP_{50} of 0.93931 and a mask $\text{mAP}_{50:95}$ of 0.59645, with mask precision and recall of 0.90623 and 0.89947, respectively. These results confirmed that the initial configuration was already strong enough to serve as a realistic benchmark rather than only as a preliminary starting point.

When evaluated on the held-out original test split at an image size of 640 pixels, the same baseline model achieved a mask mAP_{50} of 0.91921 and a mask $\text{mAP}_{50:95}$ of 0.58295. This test result is slightly lower than the best validation value, which is expected, but it still indicates good generalization to unseen elevator-panel images.

4.2.2 Hyperparameter tuning results

The three-stage hyperparameter tuning procedure improved validation performance, but the gain was moderate. Stage 1 acted as a narrow search around the baseline recipe. Only a small subset of trials produced stable non-zero mask performance, and only these candidates were carried forward. The corresponding Stage 1 validation curves are provided in Appendix A.1.

Figure 4.1 shows the broader Stage 2 validation mask $\text{mAP}_{50:95}$ trends for the successful unique Stage 2 configurations. Eight out of sixteen unique Stage 2 configurations achieved non-zero mask $\text{mAP}_{50:95}$ values. The strongest Stage 2 result was obtained with batch size 16 and image size 768, which was then forwarded to Stage 3 refinement.

Table 4.2 summarizes the best validation results of the baseline model and the best Stage 3 tuned model. The tuned model improved mask $\text{mAP}_{50:95}$ from 0.59645 to 0.63064 on validation, but this must be interpreted as a model-selection result rather than as final evidence of better generalization.

Figure 4.2 shows the corrected Stage 3 comparison between refinement from the strongest Stage 2 configuration (batch size 16, image size 768) and refinement from the second-best unique Stage 2 configuration (batch size 16, image size 640). The 768-pixel branch remains consistently stronger throughout training, reaching approximately 0.63 mask $\text{mAP}_{50:95}$, whereas the 640-pixel branch stabilizes near 0.56. The corresponding validation segmentation-loss comparison is provided in Appendix A.1. Since both runs followed the same Stage 3 refinement schedule, the




Figure 4.1: Validation mask $mAP_{50:95}$ for the successful unique Stage 2 capacity-tuning configurations.

higher accuracy of the 768-pixel branch is attributed to the inherited higher-resolution Stage 2 configuration.




Figure 4.2: Stage 3 validation mask $mAP_{50:95}$ for the two refinement branches.

The held-out test comparison in Table 4.3 does not support replacing the baseline checkpoint with the Stage 3 model. The tuning procedure improved validation performance, but the baseline remained better on the final test set and was therefore retained as the reference model for later deployment-oriented experiments within the YOLO family. A supplementary cross-architecture comparison against the final Mask R-CNN checkpoint is reported in Appendix B.2.

4.2.3 Comparison with YOLO26

YOLO26-small was evaluated under comparable training conditions to check whether the newer architecture could outperform the YOLOv8s segmentation baseline. As shown in Table 4.4, it did not. The longer 180-epoch YOLO26 run improved its validation curve slightly, but the

Table 4.2: Best validation results of the baseline and tuned YOLOv8s segmentation models.

Model	Best epoch	Mask mAP ₅₀	Mask mAP _{50:95}	Mask precision	Mask recall
Baseline YOLOv8s segmentation	49	0.93931	0.59645	0.90623	0.89947
Stage 3 tuned YOLOv8s segmentation	120	0.94268	0.63064	0.90150	0.90810

Table 4.3: Held-out test-set comparison between the baseline and the selected Stage 3 tuned YOLOv8s segmentation model.

Model	Image size	Mask mAP ₅₀	Mask mAP _{50:95}	Mask precision	Mask recall
Baseline YOLOv8s segmentation	640	0.91921	0.58295	0.86835	0.89896
Stage 3 tuned YOLOv8s segmentation	640	0.91075	0.55729	0.86537	0.87685
Baseline YOLOv8s segmentation	768	0.92384	0.62325	0.86577	0.91939
Stage 3 tuned YOLOv8s segmentation	768	0.92479	0.60845	0.87224	0.90581

final test-set ranking remained in favor of the baseline model. The corresponding validation comparison is included in Appendix A.1.

Table 4.4: Test-set comparison between the baseline YOLOv8s segmentation model and YOLO26-small.

Model	Image size	Mask mAP ₅₀	Mask mAP _{50:95}	Mask precision	Mask recall
YOLOv8s segmentation baseline	640	0.91921	0.58295	0.86835	0.89896
YOLO26-small	640	0.91209	0.57538	0.87300	0.89879
YOLOv8s segmentation baseline	768	0.92384	0.62325	0.86577	0.91939
YOLO26-small	768	0.92596	0.61689	0.88408	0.91326

4.2.4 Environment-specific fine-tuning with ADC data

The final YOLO experiment addressed domain adaptation to the target university environment. The baseline checkpoint was fine-tuned on the original dataset augmented with 46 newly collected and annotated images from HAN University R29. These images were collected and annotated by the ADC and were then manually reviewed and adjusted where needed before training. These images were split into 32 training images, 7 validation images, and 7 held-out environment-specific test images. Table 4.5 shows that this adaptation step did not materially degrade performance on the original held-out test split: mask mAP_{50:95} remained effectively unchanged (0.58295 vs. 0.58293). The validation curve in Figure 4.3 is consistent with this observation.

The effect on the environment-specific test set is much stronger, as shown in Table 4.6. Mask mAP_{50:95} increased from 0.12290 to 0.50463 after fine-tuning, while precision and recall also improved substantially. The low starting value of 0.12290 is consistent with strong domain shift: before adaptation, the baseline model had never seen the HAN University R29 environment, whose panel appearance, image framing, lighting conditions, and camera-to-panel distance differed from the original training distribution. Because this environment-specific test split contained only seven fully unseen images, even a small number of poor predictions had a strong effect on the aggregate score. This is the most practically important YOLO result of the study: a relatively small amount of deployment-specific data produced a large gain in the target

Table 4.5: Performance on the original held-out elevator test split before and after ADC fine-tuning.

Model	Mask mAP ₅₀	Mask mAP _{50:95}	Mask precision	Mask recall
Baseline YOLOv8s segmentation	0.91921	0.58295	0.86835	0.89896
ADC fine-tuned YOLOv8s segmentation	0.91912	0.58293	0.88043	0.90050




Figure 4.3: Validation mask mAP_{50:95} during ADC-based fine-tuning.

environment without sacrificing the original-domain test performance.

Table 4.6: Performance on the environment-specific ADC test split before and after ADC fine-tuning.

Model	Mask mAP ₅₀	Mask mAP _{50:95}	Mask precision	Mask recall
Baseline YOLOv8s segmentation	0.53520	0.12290	0.61209	0.70136
ADC fine-tuned YOLOv8s segmentation	0.99500	0.50463	1.00000	0.99712

To complement the held-out image-based evaluation, a controlled single-button experiment was performed on a ROS bag recorded in HAN Building R29. During this recording, the mobile cart was moved relative to one outside panel containing only a single true UP button, so both detector versions could be evaluated on the same raw image stream under controlled conditions. Table 4.7 reports the resulting filtered comparison, while Figure 4.4 shows the same experiment as confidence versus estimated distance. The plot makes the practical effect of the ADC fine-tuning visible: the fine-tuned model maintains denser detections toward the far end of the observed distance range. Moreover, a clear trend of confidence drop is observed as distance grows.

Manual review of the baseline-model outputs showed that all twelve detections above 1.5 m were false positives. These detections were therefore excluded from the table, because otherwise the baseline model would appear to gain valid recognition distance from long-range hallucinations. After this correction, the confidence difference between the two models was small, but the ADC-fine-tuned model still produced more detected frames and longer mean and median valid-detection distances. Figure 4.4 reinforces this interpretation, because the baseline

Table 4.7: Controlled single-button replay comparison on the same recorded ROS bag from HAN Building R29.

Metric	Baseline model	ADC fine-tuned model
Frames with detected buttons	106	138
Mean target confidence	0.8074	0.8069
Median target confidence	0.8444	0.8423
Mean valid-target distance [m]	0.3527	0.4282
Median valid-target distance [m]	0.3308	0.4102
Maximum valid-target distance [m]	0.8761	0.8761




Figure 4.4: Controlled single-button replay result: confidence versus estimated distance for the baseline and ADC-fine-tuned YOLO models.

detections become visibly sparser as distance increases, whereas the ADC-fine-tuned model continues to produce detections more consistently. The controlled single-button experiment therefore supports the same practical conclusion as the image-based environment-specific test set: deployment-specific R29 data improved the robustness of the detector in the target environment.

Appendix C reports two supplementary deployment analyses: the unfiltered single-button detector comparison, which makes the baseline model’s long-range false-positive behavior explicit, and the combined YOLO and OCR view of the same controlled run.

Overall, the YOLO results show three main findings. First, the baseline YOLOv8s segmentation model was already strong and difficult to surpass on the general elevator test set. Second, staged hyperparameter tuning improved validation performance but did not fundamentally change the final model ranking. Third, the most practically valuable improvement was achieved through environment-specific fine-tuning with ADC data, which substantially increased performance on the target deployment domain while preserving performance on the original dataset.

4.3 OCR hyperparameter tuning and evaluation

The OCR experiments focused on improving cropped-button recognition while keeping the recognizer lightweight enough for later embedded deployment. As in the YOLO experiments, model selection during OCR tuning was based on the validation split, and the final comparison between OCR candidates was performed on the held-out OCR test set.

4.3.1 OCR hyperparameter tuning

Stage 1: epoch selection

The first OCR tuning stage varied only the number of training epochs while keeping the optimizer configuration fixed. The main objective of this stage was to determine whether the original 60-epoch fine-tuning recipe was necessary or whether a shorter training budget would already achieve the same validation quality.

Table 4.8 shows that the 50-epoch and 60-epoch runs converged to almost identical validation performance, both reaching their best result at epoch 46. In contrast, the 30-epoch run underperformed clearly. This indicates that the recognizer required more than 30 epochs to converge, but that training beyond 50 epochs provided no meaningful additional benefit. For that reason, 50 epochs were selected as the fixed budget for the next stage of OCR hyperparameter tuning. This interpretation is also supported by the Stage 1 training-loss curves shown in Figure 4.5, where the main decrease occurs well before the end of the 50-epoch run and the longer schedule does not suggest a qualitatively different convergence regime.

Stage 2: learning rate and weight decay sweep

After fixing the training budget at 50 epochs, a second tuning stage evaluated a narrow optimizer sweep over learning rate and weight decay. Four configurations were tested. The resulting validation scores are summarized in Table 4.9.




Figure 4.5: Stage 1 OCR training loss for the epoch-selection runs.

Table 4.8: Stage 1 OCR epoch-sweep results on the validation split.

Trial	Best epoch	Accuracy	Normalized edit distance
30 epochs	16	0.78776	0.83732
50 epochs	46	0.84124	0.87080
60 epochs	46	0.84057	0.86893

The best configuration used an initial learning rate of 3×10^{-4} and a weight decay of 3×10^{-5} . This run achieved a validation accuracy of 0.84157 and a normalized edit distance of 0.87098, making it the strongest OCR candidate of the sweep. The configuration with the same learning rate but lower weight decay performed substantially worse, while the run using 5×10^{-4} degraded sharply and converged to a much lower performance level. Increasing the learning rate further to 7×10^{-4} recovered most of the performance, but still remained below the best 3×10^{-4} run.

These results indicate that OCR fine-tuning was sensitive to the optimizer settings even within a relatively narrow search range. In particular, the stage-2 results suggest that a slightly lower learning rate than the original baseline provided more stable and ultimately better convergence on the recognition task. This behavior is visible in Figure 4.6, where the best run maintains the highest evaluation accuracy. The corresponding normalized edit-distance curves are provided in Appendix A.2.

Table 4.9: Stage 2 OCR hyperparameter sweep on the validation split.

Trial	Learning rate	Weight decay	Accuracy	Normalized edit distance
lr3e4.wd1e5.ep50	3×10^{-4}	1×10^{-5}	0.77531	0.83148
lr3e4.wd3e5.ep50	3×10^{-4}	3×10^{-5}	0.84157	0.87098
lr5e4.wd3e5.ep50	5×10^{-4}	3×10^{-5}	0.59603	0.74899
lr7e4.wd3e5.ep50	7×10^{-4}	3×10^{-5}	0.83653	0.86546




Figure 4.6: Stage 2 OCR validation accuracy for the learning-rate and weight-decay sweep.

4.3.2 Final OCR benchmark

The final OCR benchmark compared the original fine-tuned recognizer against the best-performing stage-2 configuration on the held-out OCR test set. The results are shown in Table 4.10. The stage-2 model outperformed the initial OCR model on both recognition-quality metrics. Test accuracy improved from 0.81635 to 0.83486, while normalized edit distance improved from 0.85574 to 0.87190. The throughput values were measured offline on the training workstation and are not directly comparable to the Jetson deployment timings reported later.

Taken together, these results show that the OCR hyperparameter tuning process produced a real improvement rather than only a validation-specific gain. Unlike the YOLO tuning experiments, where better validation performance did not translate into a better final test-set model, the OCR stage-2 winner also improved the held-out test benchmark. For this reason, the stage-2 best OCR model was selected as the final recognition candidate for later integration into the complete perception pipeline.

Table 4.10: Offline held-out OCR test-set comparison on the training workstation between the initial recognizer and the selected stage-2 OCR model.

Model	Accuracy	Normalized edit distance	Throughput (samples/s)
Initial OCR model	0.81635	0.85574	805.68
Stage-2 best OCR model	0.83486	0.87190	817.71

4.4 Model hardware optimization results

As described in Chapter 3, the final YOLO and OCR models were exported to TensorRT for optimized execution on the NVIDIA Jetson platform. To verify the effect of this export step, a compact benchmark was performed on the held-out test image `367.jpg`, which is included in Appendix B.3. For YOLO, the comparison used the PyTorch checkpoint and the exported TensorRT engine on the same full image. For OCR, the comparison used the Paddle inference

model and the exported TensorRT engine on the five annotated button crops from the same image.

Table 4.11: Illustrative single-image hardware optimization micro-benchmark on test image 367.jpg.

Model stage	Baseline backend	TensorRT backend	Speedup
YOLO inference [ms/image]	30.29	17.86	1.70×
YOLO wall time [ms/image]	44.66	38.98	1.15×
OCR wall time [ms/crop]	167.72	10.29	16.30×

The updated YOLO micro-benchmark compares the original PyTorch checkpoint against the final deployment-oriented TensorRT engine of the ADC-fine-tuned model. On the same held-out test image, TensorRT reduced the mean measured YOLO inference time from 30.29 ms/image to 17.86 ms/image, while the corresponding single-image wall time also improved from 44.66 ms/image to 38.98 ms/image. The OCR result is more direct: the TensorRT engine reduced average per-crop wall time from 167.72 ms to 10.29 ms while preserving the same accepted predictions on all five button regions in the test image. As a single-image micro-benchmark, this result is illustrative rather than a substitute for the broader ROS 2 pipeline timing measurements reported later in this chapter. It nevertheless supports the use of TensorRT export as an important deployment optimization step before integrating the models into the full ROS pipeline.

4.5 ROS pipeline performance

After the final YOLO and OCR models were exported to TensorRT, the complete perception stack was evaluated on the NVIDIA Jetson Orin Nano in its deployed ROS configuration. The purpose of these experiments was not to compare neural-network accuracy, but to quantify the practical throughput of the integrated pipeline under different runtime settings. During these experiments, the Stereolabs ZED X camera was launched with the `zedx` wrapper profile, which uses a native HD1200 grab mode at 30 Hz and publishes lens-corrected color and depth data at 15 Frames Per Second (FPS). With the configured downscale factor of 2.0, the effective published image resolution used by the perception pipeline was 960×600 pixels.

Three final runtime scenarios were compared:

- camera + YOLO + OCR + ADC,
- camera + YOLO + OCR without collector and with OCR uncapped, and
- camera + YOLO + OCR without collector and with OCR limited to four regions per message.

The OCR input was capped at four regions per message as a pragmatic bound to limit recognition workload and improve end-to-end runtime.

Table 4.12 summarizes the resulting topic throughput. The collector-enabled configuration sustained approximately 10 FPS while preserving the full deployment data-collection path. When the collector was disabled and OCR was allowed to process all candidate regions, the

end-to-end rate increased to about 12.2 FPS. The best runtime-oriented configuration was obtained when OCR was capped to four regions per message, yielding approximately 13.1 FPS end-to-end on the final `/button.targets_3d_ocr` output topic. The reported topic rates were measured with the standard `ros2 topic hz` utility on the relevant output topics.

Table 4.12: Final ROS pipeline throughput on the NVIDIA Jetson Orin Nano after the ROI-mask refactor, measured with `ros2 topic hz`.

Configuration	<code>/button.targets_3d</code> FPS	<code>/button.targets_3d_ocr</code> FPS
With collector	~ 10.0	~ 10.0
No collector, OCR uncapped	~ 12.15	~ 12.29
No collector, OCR capped to 4 ROIs	~ 13.14–13.47	~ 13.08–13.25

The timing breakdown confirms that the final system bottleneck is not the TensorRT forward pass alone, but the combined cost of segmentation postprocessing, depth estimation, and OCR execution inside the live ROS graph. Table 4.13 shows typical timing ranges of the YOLO and OCR nodes for the three runtime configurations. A more detailed stage-level breakdown of the final runtime configuration is provided in Appendix E.4. In collector mode, the YOLO node required approximately 81–91 ms per callback and the OCR node 79–95 ms. Removing the collector while keeping OCR uncapped reduced the practical throughput penalty of data export but still required roughly 11 processed OCR regions per message, which kept OCR total time in the 77–102 ms range. Limiting OCR to four regions reduced OCR total time to roughly 33–35 ms and allowed the complete pipeline to reach its best measured throughput.

Table 4.13: Typical node timing ranges for the final ROS runtime experiments.

Configuration	YOLO total [ms]	YOLO depth [ms]	OCR total [ms]	OCR processed ROIs
With collector	~ 81–91	~ 3.6–6.6	~ 79–95	not capped
No collector, OCR uncapped	~ 78–99	~ 4–7	~ 77–102	~ 11
No collector, OCR capped to 4 ROIs	~ 70.4–77.8	~ 5.3–7.0	~ 33.0–35.2	4.0

These results show that the integrated perception pipeline can operate at deployment-relevant frame rates on the NVIDIA Jetson Orin Nano, provided that annotation serialization is kept out of the real-time YOLO path and OCR workload is bounded. For runtime operation, the most effective configuration is the ROI-limited OCR setup without concurrent collection, which achieved approximately 13.1 FPS. For deployment sessions that must simultaneously gather new training data, the collector-enabled configuration remains practical at approximately 10 FPS, which is sufficient to preserve the full adaptive data-collection loop described in Chapter 3.

4.5.1 Qualitative end-to-end pipeline output

To simulate realistic environment test, a longer and more diverse university elevator ROS bag was therefore collected and replayed and its detections were manually reviewed as true positives or false positives. Figure 4.8 shows the reviewed YOLO detection confidence versus estimated target distance for that deployment-style run. The plot makes explicit that high-confidence accepted detections are concentrated at shorter distances, while lower-confidence and more failure-prone cases occur as the target becomes more distant in the realistic university environment. A fuller supporting appendix analysis is provided in Appendix C.

Figure 4.7 then illustrates the final deployed output of the integrated perception pipeline for an in-cabin elevator panel observed during the live testing with the mobile cart in Han Building R29. These overlays are shown as qualitative examples of the final `/button_targets_3d_ocr` output rather than as the source of the controlled single-button replay comparison from Section 4.2.4. A fuller appendix example of the same type of final output, including the full `/button_targets_3d_ocr` message, is provided in Appendix E.5. Both overlays are derived from the same accepted `/button_targets_3d_ocr` message set, but visualize different aspects of the final result. The first view emphasizes the segmentation masks, representative button centers, and estimated distance values, while the second view emphasizes bounding boxes, recognized labels, and the corresponding YOLO confidence scores.



Figure 4.7: Representative end-to-end output of the final deployed ROS perception pipeline for an in-cabin elevator panel.




Figure 4.8: Reviewed YOLO detection confidence versus estimated distance on a realistic university elevator replay run. True and false detections were manually reviewed after replay.

This example makes the final system output more explicit than throughput tables alone. Instead of exposing only intermediate detections, the deployed pipeline produces a task-oriented target representation that already combines 2D button localization, a representative image center for depth projection, semantic text interpretation, and an estimated spatial distance for downstream robot use.

5 Discussion

This chapter interprets the results presented in Chapter 4 and relates them to the project objective of developing a scalable computer-vision workflow for robot interaction with elevator infrastructure on embedded hardware. Taken together, the implemented stages instantiate the main loop outlined earlier in the workflow-and-architecture overview in Figure 3.1: data collection, annotation support, retraining, deployment, and environment-specific refinement.

5.1 Interpretation of the main findings

The main outcome of this project is that a practical embedded perception workflow can be achieved by combining instance-segmentation-based button detection, crop-based OCR, and a modular ROS runtime pipeline. The results show that the strongest overall system did not come from selecting the newest model architecture or from maximizing validation performance in isolation. Instead, the most effective workflow emerged from a sequence of engineering decisions that balanced perception quality, deployment constraints, and future scalability.

For YOLO, the baseline YOLOv8s segmentation model remained the most reliable choice for deployment. Hyperparameter tuning improved validation performance, but those gains did not translate into a better final held-out test model. A plausible explanation is that the tuning process optimized more strongly for the validation split, while the baseline configuration generalized slightly better to unseen test images. Likewise, YOLO26-small did not outperform the baseline under the same evaluation conditions. This is an important result because it shows that, for this task, a carefully chosen and well-understood baseline can be more valuable than adopting a newer architecture without a clear deployment benefit. The most practically useful YOLO improvement came from environment-specific fine-tuning with ADC data from Han Building R29, which substantially improved performance in the target environment while preserving the original-domain performance.

For OCR, the project produced a clearer tuning outcome. The selected stage-2 configuration improved both test accuracy and normalized edit distance relative to the initial recognizer. In contrast to the YOLO tuning results, the OCR validation improvements did translate into better held-out test-set performance. This suggests that the OCR stage was more sensitive to optimization settings in a way that remained meaningful at test time.

At system level, the deployed ROS pipeline demonstrates that model quality alone is not sufficient for real-time embedded performance. The final pipeline throughput depended strongly on how much non-inference work remained in the runtime path. The shift from full-image mask export to compact ROI-local mask handling, together with the split-node YOLO/OCR architecture and the bounded OCR workload, produced the most important runtime gains. The best measured configuration achieved about 13.1 FPS, while the collector-enabled configuration remained practical at about 10 FPS. This shows that the embedded performance target was not reached by model export alone, but by a combination of model export, message-design decisions, and workload placement.

5.2 Relation to the research questions

The results provide a coherent answer to the main research question. A computer-vision workflow for hospital elevator-button perception can be designed by using instance segmentation for button localization, crop-based OCR for semantic interpretation, TensorRT export for embedded execution, and an ROS-based split-node architecture for modular runtime integration.

The first sub-question concerned the relevant human-interaction points in hospital environments. The project focused on elevator buttons and related panel interfaces because they are a concrete and operationally critical interaction class for multi-floor service-robot mobility. This focus was also aligned with Ambyon’s immediate deployment priorities, which made elevator-button perception relevant and high value interaction target for the present project. This focus proved appropriate: the task is sufficiently challenging to expose problems of detection, segmentation, recognition, and embedded deployment, while still being narrow enough to support a complete workflow.

The second sub-question asked which methods and architectures were suitable. The literature review suggested that instance segmentation is better aligned with precise robot interaction than coarse object detection alone, and the results support this choice. The segmentation masks were directly useful for masked depth estimation and 3D target generation. The final architecture therefore combined a segmentation-based detector with a separate recognizer rather than relying on a single end-to-end model.

The third sub-question concerned dataset design. The results show that dataset quality and dataset relevance were both important. The original curated dataset was sufficient to train a strong baseline, but the environment-specific fine-tuning experiment demonstrated that even a relatively small amount of deployment-collected data can substantially improve performance in the target domain. In the controlled R29 single-button experiment, the adapted detector produced more frames with detections and longer typical valid-detection distances once the baseline model’s manually reviewed long-range false positives were excluded from the valid-distance comparison. The reviewed deployment-style university run further showed that accepted detection confidence degrades as distance increases. This supports the decision to include an ADC in the full workflow and suggests that further fine-tuning on carefully collected deployment data is likely to remain a high-value improvement path.

The fourth sub-question asked which optimization techniques support embedded inference. Here the answer is twofold. First, TensorRT export was beneficial for both YOLO and OCR, with especially strong gains for OCR. Second, export alone was not enough: runtime-side optimizations such as reducing annotation overhead, separating the inference nodes, and limiting the number of processed OCR regions were equally important.

The fifth sub-question concerned scalability and efficient future improvement. The final system addresses this directly through the ADC, which turns runtime use into a source of curated retraining data. The environment-specific YOLO improvement demonstrates that this design is not only theoretically scalable, but also practically useful.

5.3 Limitations

Several limitations should be considered when interpreting the results. First, the project focused on elevator-interface perception rather than the full set of hospital interaction elements. While this was a justified and useful scope reduction, the final conclusions should not automatically be generalized to all door panels, access readers, or human-machine interfaces.

Second, some benchmark results remain sensitive to the evaluation setup. For example, the single-image framework-to-TensorRT comparison is useful for illustrating acceleration trends, but it is not by itself a complete deployment benchmark. Likewise, some distance-related YOLO improvements were most clearly observed in field behavior and specific replay analyses rather than in a single universal summary metric.

Third, the project was executed on a single embedded platform and camera setup. The NVIDIA Jetson Orin Nano and Stereolabs ZED X are representative for the target use case, but performance and tuning decisions may differ on other edge devices or other sensing configurations. In addition, the current Jetson setup still runs from a microSD card rather than an SSD, which is not ideal for long logging sessions, dataset export, or sustained deployment experiments.

Fourth, the report currently emphasizes perception and runtime integration rather than downstream manipulation or actuation. The system produces 3D button targets and button labels, but the full closed-loop robot interaction behavior lies beyond the present project scope.

Fifth, the current ADC closes the loop only partially. It can collect, filter, and store samples in a format that is already useful for later YOLO and OCR updates, but there is still a human in the loop for verification. In addition, export from the robot, transfer of the dataset session, integration with the CVAT tooling, and redeployment of updated models are not yet fully automated. The improvement cycle is therefore structured, but not yet fully closed end to end.

5.4 Implications for future robot deployment

Despite these limitations, the project has clear practical implications. The results suggest that a service robot does not require a monolithic perception stack to achieve useful embedded performance. A modular pipeline with explicit message interfaces can be easier to tune, benchmark, and extend. This is particularly important for educational and research robots such as Care-BOTS, where iterative development and subsystem replacement are expected.

A further practical deployment consideration is model licensing. The YOLO framework used in this project is distributed by Ultralytics under the AGPL-3.0 open-source license, with separate commercial licensing options for enterprise use. In contrast, PaddleOCR and Torchvision are distributed under the more permissive Apache 2.0 and BSD-3-Clause licenses, respectively. Although licensing did not affect the technical validity of the experiments reported here, it should be considered in any future transition from research prototyping to commercial or operational deployment.

The project also suggests that deployment data should be treated as a core part of the system design rather than as an afterthought. The inclusion of the ADC makes the perception workflow self-improving in a structured way: the same runtime system used in deployment can

collect higher-value examples for the next training cycle. This closes an important loop between experimentation, deployment, and future model improvement.

6 Conclusions and recommendations

This chapter summarizes the final conclusions of the project and translates the main findings into concrete recommendations for future work.

6.1 Conclusions

The objective of this project was to develop and validate an efficient and scalable computer-vision workflow that enables a service-robot platform to detect and classify human-interaction elements in hospital environments using embedded hardware. Based on the results presented in this report, this objective was achieved for the targeted elevator-interaction use case.

The project demonstrates that a practical embedded perception workflow can be constructed from four main components: a strong instance-segmentation baseline for button localization, a crop-based OCR stage for button interpretation, TensorRT export for hardware-efficient inference, and a modular ROS pipeline that integrates detection, 3D projection, text recognition, and deployment data collection.

The main conclusions are as follows.

1. Instance segmentation is an effective method for elevator-button perception because it supports both accurate 2D localization and mask-based depth estimation for 3D target generation.
2. The baseline YOLOv8s segmentation model was the strongest deployment candidate among the evaluated detector variants. Hyperparameter tuning improved validation performance, but not the final held-out test ranking.
3. Environment-specific fine-tuning with ADC data from Han Building R29 produced the most practically valuable YOLO improvement, substantially increasing target-domain performance while preserving original-domain performance. In the controlled single-button R29 experiment, the adapted detector produced more frames with detections (138 vs. 106) and longer mean and median valid-detection distances (0.4282 m vs. 0.3527 m and 0.4102 m vs. 0.3308 m, respectively) after the initial model’s manually reviewed long-range false positives were excluded from the valid-distance comparison. The corresponding confidence values were very similar, indicating that the main practical gain was more reliable target coverage rather than a large confidence increase. A separate reviewed deployment-style university run further showed a clear practical trend: accepted YOLO detection confidence decreases as distance increases, which strengthens the case for treating distance as an important reliability cue during deployment.
4. The selected stage-2 OCR model improved held-out recognition quality and formed a suitable semantic interpretation stage for the full pipeline.
5. TensorRT export was beneficial for both deployed models. The effect was moderate but clear for YOLO inference time, and very strong for OCR per-crop wall time.

6. The final ROS pipeline achieved practical embedded runtime performance on the NVIDIA Jetson Orin Nano, with the best measured configuration operating at about 13.1 FPS (approximately 76 ms/frame) and the collector-enabled configuration remaining usable at about 10 FPS (approximately 100 ms/frame).
7. Scalability is supported not only by model architecture choices, but also by the inclusion of the ADC, which enables structured deployment-time data collection for future retraining.

Taken together, these conclusions answer the main research question: a practical computer-vision workflow for elevator-button perception can be achieved on embedded hardware by combining segmentation-based detection, task-specific recognition, deployment-oriented optimization, and iterative deployment data collection within a modular ROS architecture.

6.2 Recommendations

The following recommendations are proposed for future work and further development of the system.

1. Extend the perception scope beyond elevator buttons to additional hospital interaction elements such as access readers, door controls, and other panel interfaces.
2. Perform larger and more controlled deployment trials in multiple buildings so that domain shift, lighting variation, and panel diversity can be evaluated more systematically.
3. Continue using the ADC as part of regular deployment, with periodic review, further fine-tuning, and retraining cycles based on newly collected field samples.
4. Strengthen the hyperparameter tuning process for both the YOLO and OCR models by using broader and more systematic searches, so that future model updates are based on more stable and better-validated configurations.
5. Automate the improvement cycle further by reducing manual steps between ADC export, annotation verification, training, and redeployment, so that future model updates require less manual coordination.
6. Perform real robot trials in hospital-like and, when possible, actual hospital environments so that the perception pipeline can be validated under realistic operational constraints and new deployment data can be collected systematically.
7. Replace the current Jetson microSD-based storage setup with an SSD-based configuration to improve robustness for logging, dataset export, and long deployment sessions.
8. Reimplement runtime-critical ROS nodes in C++ where appropriate, since the current Python implementation is practical for development but likely leaves performance headroom on the embedded platform.
9. Evaluate the pipeline on additional embedded platforms to determine how transferable the optimization choices are beyond the current NVIDIA Jetson-based setup.

10. Integrate the perception pipeline with downstream robot action modules so that the 3D target estimates can be validated in full interaction tasks rather than perception-only experiments.

6.3 Final reflection

This project provides a concrete foundation for autonomous interaction with elevator infrastructure in hospital-like environments. More broadly, it shows that robust embedded perception is not achieved by model selection alone. The best results emerged from treating dataset design, model choice, runtime architecture, and deployment data collection as one connected system. That systems perspective is likely to remain essential as the CareBOTS platform evolves toward more complete autonomous operation in real healthcare environments.

Bibliography

- [1] Bai, J., Lu, F., Zhang, K., et al. (2019). Onnx: Open neural network exchange. In *Proceedings of the 2019 ACM International Conference on Multimedia Systems*, pages 460–466.
- [2] Bernard, A., Hörnle, T., and Dillmann, R. (2020). Vision-based perception of control interfaces for robotic manipulation. *Robotics and Autonomous Systems*, 124:103386.
- [3] Chen, T., Chen, Z., and Chen, Y. (2018). Deep learning in robotics: A review of recent research. *Advanced Robotics*, 32(5):201–214.
- [4] Cordts, M., Omran, M., Ramos, S., Rehfeld, T., Enzweiler, M., Benenson, R., Franke, U., Roth, S., and Schiele, B. (2016). The cityscapes dataset for semantic urban scene understanding. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3213–3223.
- [5] He, K., Gkioxari, G., Dollár, P., and Girshick, R. (2017). Mask r-cnn. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2961–2969.
- [6] Holland, J. et al. (2021). Service robots in the healthcare sector. *Robotics*, 10(1):47.
- [7] Iguazio (2020). Nuclio: High-performance serverless event and data processing platform. <https://nuclio.io>.
- [8] Jocher, G., Chaurasia, A., and Qiu, J. (2023). YOLOv8: Ultralytics next-generation yolo models. *arXiv preprint arXiv:2305.09972*.
- [9] Kang, Y., Kim, H., and Lee, J. (2020). Edge ai: On-device intelligence for embedded systems. *IEEE Consumer Electronics Magazine*, 9(4):20–27.
- [10] Lin, T.-Y., Maire, M., Belongie, S., et al. (2014). Microsoft coco: Common objects in context. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 740–755.
- [11] Liu, H. (2022). Indoor scene understanding for the visually impaired based on semantic segmentation. Thesis. Figure 3.1 accessed via ResearchGate on 21 March 2026.
- [12] Liu, J., Fang, Y., Zhu, D., Ma, N., Pan, J., and Meng, M. Q.-H. (2021). A large-scale dataset for benchmarking elevator button segmentation and character recognition. *IEEE Transactions on Industrial Informatics*.
- [13] Long, J., Shelhamer, E., and Darrell, T. (2015). Fully convolutional networks for semantic segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3431–3440.
- [14] Martínez, M., Montero, R., and Pérez, E. (2018). Irso: A dataset for interactive robotic service operations in indoor environments. In *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1234–1241.

- [15] Masala, G. L. et al. (2025). Artificial intelligence and assistive robotics in healthcare: A narrative review. *International Journal of Environmental Research and Public Health*, 22(5):781.
- [16] Merkel, D. (2014). Docker: Lightweight linux containers for consistent development and deployment. *Linux Journal*, 2014(239).
- [17] PaddlePaddle Authors (2026). Paddleocr documentation. <https://paddlepaddle.github.io/PaddleOCR/main/en/index.html>. Accessed 21 March 2026.
- [18] Redmon, J., Divvala, S., Girshick, R., and Farhadi, A. (2016). You only look once: Unified, real-time object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788.
- [19] Ren, S., He, K., Girshick, R., and Sun, J. (2017). Faster r-cnn: Towards real-time object detection with region proposal networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(6):1137–1149.
- [20] Sekachev, B., Manovich, N., Zhiltsov, M., et al. (2020). Cvat: Computer vision annotation tool. *arXiv preprint arXiv:2009.13245*.
- [21] Sun, X., Wu, J., Zhang, X., Zhang, Z., Zhang, C., Xue, T., Freeman, W., and Tenenbaum, J. (2017). Scenenet rgb-d: Can 5m synthetic images beat generic imagenet pre-training? In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2678–2687.
- [22] Tan, M. and Le, Q. (2019). Efficientnet: Rethinking model scaling for convolutional neural networks. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 6105–6114.
- [23] Ultralytics (2026). Ultralytics yolo documentation. <https://docs.ultralytics.com/>. Accessed 21 March 2026.
- [24] Wang, X., Luo, C., and Zhang, Y. (2019). Vision-based perception for autonomous indoor service robots: A survey. *IEEE Transactions on Automation Science and Engineering*, 16(4):1603–1619.

A Supplementary hyperparameter tuning curves

A.1 YOLO hyperparameter tuning curves

This appendix collects supporting YOLO hyperparameter-tuning figures that were removed from the main results chapter to keep the core model-selection narrative concise.




Figure A.1: Validation mask mAP_{50:95} for the successful Stage 1 hyperparameter trials.




Figure A.2: Stage 3 validation segmentation loss for the two refinement branches.




Figure A.3: Validation mask $mAP_{50:95}$ for the two YOLO26 training runs.

A.2 OCR hyperparameter tuning curves

This appendix collects supporting OCR hyperparameter-tuning curves that complement the main validation-accuracy figure and the stage-summary tables in Chapter 4.




Figure A.4: Stage 2 OCR normalized edit distance for the learning-rate and weight-decay sweep.

B Supplementary model evaluations

B.1 Mask R-CNN checkpoint progression during annotation support

The Mask R-CNN model was used as an annotation-support model within the CVAT–Nuclio workflow, rather than as the final deployment model. During dataset construction, the model was retrained several times as more annotated images became available. To quantify the effect of this iterative retraining, four checkpoints were evaluated on the same held-out test split used for the annotation dataset. Table B.1 reports both bounding-box and mask metrics, with mask AP treated as the primary segmentation indicator.

Table B.1: Held-out test-set performance of the annotation-support Mask R-CNN checkpoints as the training set grew during the annotation process.

Checkpoint stage	Training images	BBox AP@[.50:.95]	Mask AP@[.50:.95]	Mask AP@0.50
Initial checkpoint	90	0.5878	0.5844	0.7776
Retrained checkpoint	449	0.6682	0.6660	0.8284
Retrained checkpoint	900	0.7007	0.6903	0.8772
Final checkpoint	1400	0.7192	0.7000	0.8990

The results show a consistent improvement trend as the amount of annotated data increased. The largest gain was observed between the initial 90-image checkpoint and the first larger retrained checkpoint at 449 images, while later retraining to 900 and 1400 images provided smaller but still measurable improvements. On the primary mask metric, AP@[0.50:0.95] increased from 0.5844 to 0.7000. This supports the annotation strategy adopted in this project: iterative retraining of the pre-annotation model improved the quality of automatically proposed instances, thereby reducing the manual correction burden in later annotation rounds.

B.2 Baseline YOLOv8s segmentation versus final Mask R-CNN

A direct cross-architecture comparison was performed between the baseline YOLOv8s segmentation checkpoint and the final 1400-image Mask R-CNN checkpoint on the original held-out elevator test split. Official mask mAP₅₀ and mAP_{50:95} values were taken from the respective test-set evaluation pipelines. In addition, a direct replay comparison was run on the same 602 test images using a confidence threshold of 0.5 and greedy mask matching at IoU 0.5 to estimate thresholded precision, recall, and average inference time per image.

Table B.2: Supplementary comparison between the baseline YOLOv8s segmentation and the final annotation-support Mask R-CNN checkpoint on the original held-out test split.

Model	Mask mAP ₅₀	Mask mAP _{50:95}	Precision*	Recall*	Time [s/image]*
Baseline YOLOv8s segmentation	0.9192	0.5829	0.8787	0.9063	0.0091
Final Mask R-CNN	0.8990	0.7000	0.7578	0.9547	0.0196

The comparison shows a split result. The final Mask R-CNN checkpoint achieved higher official mask AP on the original test split, whereas the baseline YOLOv8s segmentation model provided substantially higher thresholded precision and roughly twice the throughput in the direct replay experiment. Mask R-CNN also retained higher thresholded recall. This supports the role division used in this project: Mask R-CNN was valuable as an offline annotation-support model, while YOLOv8s segmentation remained the preferred deployment-oriented detector due to its runtime efficiency and simpler integration into the embedded perception pipeline.

*Precision, recall, time, and FPS come from the direct replay script on identical test images at confidence threshold 0.5 and greedy IoU matching at 0.5; they are therefore supplementary to the official AP metrics rather than a replacement for them.

B.3 Reference image for the hardware optimization benchmark

Figure B.1 shows the held-out test image used for the single-image hardware optimization micro-benchmark in Section 4.11. It is included here so that readers can inspect the elevator-panel scene on which the PyTorch/Paddle and TensorRT inference times were compared.



Figure B.1: Held-out test image 367.jpg used in the hardware optimization micro-benchmark.

C Supplementary deployment experiments

This appendix collects supplementary deployment-oriented analyses that support the main-body results of Section 4.2.4 and Section 4.5. The main chapter presents the most compact figures under page constraints, while the appendix retains additional deployment behavior that was informative during interpretation.

C.1 Unfiltered single-button detector comparison

Table C.1 shows the same controlled R29 single-button experiment before the baseline-model detections above 1.5 m were removed. It is included to document the full deployment behavior of both models. The apparent long-range advantage of the baseline model in the unfiltered distance statistics should not be interpreted as better valid detection range, because manual review showed that all twelve baseline detections above 1.5 m were false positives.

Table C.1: Unfiltered controlled single-button comparison on the same recorded ROS bag from HAN Building R29.

Metric	Baseline model	ADC fine-tuned model
Frames with detected buttons	116	138
Mean target confidence	0.7855	0.8069
Median target confidence	0.8328	0.8423
Mean valid-target distance [m]	0.4408	0.4282
Median valid-target distance [m]	0.3353	0.4102
Maximum valid-target distance [m]	1.6536	0.8761

C.2 Single-button end-to-end confidence behavior

Besides the detector-only baseline-versus-ADC comparison shown in the main body, the controlled single-button R29 replay was also inspected at the full YOLO+OCR pipeline level. Figure C.1 shows the detector confidence together with the downstream OCR confidence for the same run. Correct UP reads remain clustered in the lower-to-mid distance range, while the few OCR errors appear at larger distances within that run.

C.3 Mixed-environment replay review by distance bin

A second replay analysis was performed on a longer university elevator bag that contained multiple panels, cluttered backgrounds, and changing viewpoints. The detections were manually reviewed as true positives or false positives after replay. Figure C.2 shows the reviewed true- and false-positive counts by distance bin for that broader run. The figure is included as supporting evidence for the realistic deployment trend discussed in Section 4.5: reviewed detections were highly reliable below roughly 1.5 m, whereas the longer-range bins were dominated by false positives.




Figure C.1: Supplementary single-button R29 replay plot showing YOLO confidence, OCR confidence, and the distances at which OCR stopped reading the UP button correctly.




Figure C.2: Reviewed true and false positive detections by distance bin on the longer university replay run. Close-range bins were sampled, while the longer-distance bins were reviewed exhaustively.

D Supplementary dataset quality plots

This appendix provides compact visual support for the dataset quality assessment summarized in Chapter 4. The figures are included only as supporting material and are not required to follow the main argument of the report.

D.1 Image-level quality distributions




Figure D.1: Distribution of the equivalent button diameter measured from the annotated instances. The x-axis shows the estimated button diameter in pixels, while the y-axis shows the number of annotated button instances in each histogram bin. This figure directly supports the button-size summary values reported in Chapter 4.

D.2 Spatial button distribution




Figure D.2: Distribution of the mean grayscale intensity across the training images. The x-axis shows the average grayscale value per image on the 0–255 intensity scale, where lower values indicate darker images and higher values indicate brighter images. The y-axis shows the number of training images in each histogram bin.




Figure D.3: Heatmap of annotated button-centroid locations in normalized image coordinates. The x-axis shows the horizontal centroid position normalized by image width, with 0 at the left image border and 1 at the right border. The y-axis shows the vertical centroid position normalized by image height, with 0 at the top of the image and 1 at the bottom. The color intensity represents how many annotated button instances fall into each two-dimensional bin.

E ROS pipeline parameter reference

This appendix summarizes the main runtime parameters of the deployed ROS perception pipeline. The tables are intended as a practical reference for reproducing experiments or adjusting node behavior during deployment. Values shown in the “Default / configured value” column correspond to the final thesis configuration where this was defined explicitly in the YAML files; otherwise the node’s declared software default is reported.

E.1 YOLO inference node

Table E.1: Main parameters of `yolo_engine_node`.

Parameter	Default / configured value	Purpose
<code>engine_path</code>	launch-time path	Path to the TensorRT YOLO engine.
<code>image_topic</code>	<code>/zed/zed_node/ rgb/color/rect/ image</code>	Input image topic.
<code>depth_topic</code>	<code>/zed/zed_node/ depth/depth\ _registered</code>	Input depth topic used for 3D projection.
<code>camera_info_topic</code>	<code>/zed/zed_node/ rgb/color/rect/ camera_info</code>	Camera intrinsics topic.
<code>detections_topic</code>	<code>/detections</code>	Optional 2D detection topic retained for debugging.
<code>targets_3d_topic</code>	<code>/button_targets\ _3d</code>	Main 3D target output topic.
<code>conf</code>	0.50	Confidence threshold for accepted detections.
<code>infer_hz</code>	8.0	Timer frequency for inference callbacks.
<code>enable_distance\ _estimation</code>	true	Enables depth-based 3D target estimation.
<code>max_depth_age_ms</code>	150.0	Maximum allowed age of buffered depth data.
<code>min_depth_m</code>	0.10	Minimum accepted depth value.
<code>max_depth_m</code>	10.0	Maximum accepted depth value.
<code>min_valid_depth\ _pixels</code>	20	Minimum number of valid masked depth pixels.
<code>depth_percentile</code>	50.0	Percentile used for representative target depth.
<code>enable_ocr</code>	false	Legacy inline OCR mode; disabled in the split-node pipeline.
<code>ocr_engine_path</code>	empty	Optional legacy inline OCR engine path.
<code>ocr_input_name</code>	x	Legacy inline OCR input tensor name.
<code>ocr_charset_path</code>	empty	Optional legacy inline OCR character set path.

Parameter	Default / configured value	Purpose
ocr_input_h	48	Legacy inline OCR input height.
ocr_input_w	320	Legacy inline OCR input width.
ocr_min_conf	0.60	Legacy inline OCR confidence threshold.
ocr_max_rois_per_ _frame	8	Legacy inline OCR ROI cap.
warmup_runs	3	Number of startup warmup inferences.
enable_mask_rle_ _export	false	Enables direct mask export for dataset collection.
publish_debug_image	true	Enables publication of a visual debug image.
debug_image_topic	/yolo/debug_ _image	Debug image output topic.
enable_perf_logging	true	Enables periodic runtime timing logs.
perf_log_every_n	50	Logging interval for performance statistics.

E.2 OCR inference node

Table E.2: Main parameters of `ocr_trt_node`.

Parameter	Default / configured value	Purpose
image_topic	/zed/zed_node/ rgb/color/rect/ image	Input image topic used for buffered crops.
targets_topic_in	/button_targets\ _3d	Input topic from the YOLO node.
targets_topic_out	/button_targets\ _3d_ocr	Output topic with text-enriched targets.
debug_image_topic	/ocr/debug_image	Debug visualization topic.
ocr_engine_path	launch-time path	Path to the TensorRT OCR engine.
ocr_input_name	x	Name of the input tensor in the engine.
ocr_charset_path	empty	Optional character set file path.
ocr_input_h	48	Input crop height for recognition.
ocr_input_w	320	Input crop width for recognition.
ocr_min_conf	0.60	Minimum accepted text confidence.
ocr_max_rois_per\ _msg	8	Maximum number of button crops processed per message.
warmup_runs	3	Number of startup warmup inferences.
max_image_age_ms	220.0	Maximum age of buffered image data matched to a target message.
image_buffer_size	30	Number of recent images retained for matching.
publish_debug_image	declared default true	Enables publication of the OCR debug image.

Parameter	Default / configured value	Purpose
<code>enable_perf_logging</code>	declared default <code>true</code>	Enables periodic timing logs.
<code>perf_log_every_n</code>	declared default 50	Logging interval for performance statistics.
<code>enable_fuzzy_map</code>	<code>true</code>	Enables post-recognition vocabulary correction.
<code>label_vocab_path</code>	configured file path	Path to known elevator button labels.
<code>fuzzy_max_dist</code>	1	Maximum edit distance for fuzzy label correction.
<code>fuzzy_min_len</code>	3	Minimum string length before fuzzy matching is applied.
<code>fuzzy_min_score</code>	0.0	Minimum score required before fuzzy correction is considered.

E.3 Active Data Collector

Table E.3: Main parameters of `active_data_collector`.

Parameter	Default / configured value	Purpose
<code>image_topic</code>	<code>/zed/zed_node/ rgb/color/rect/ image</code>	Input image topic.
<code>targets_topic</code>	<code>/button_targets\ _3d_ocr</code>	Input target topic from the OCR node.
<code>dataset_root</code>	configured dataset path	Root folder for collected sessions.
<code>session_name</code>	empty	Optional explicit session name.
<code>save_overlay_image</code>	<code>true</code>	Saves an annotated overlay image for each accepted sample.
<code>overlay_font_scale</code>	0.30	Font size used in overlay rendering.
<code>overlay_font_thickness</code>	1	Font thickness used in overlay rendering.
<code>max_image_age_ms</code>	250.0	Maximum age of the buffered image matched to a target message.
<code>image_buffer_size</code>	40	Number of recent images retained for timestamp matching.
<code>enable_perf_logging</code>	declared default <code>true</code>	Enables periodic collector timing logs.
<code>perf_log_every_n</code>	declared default 20	Logging interval for collector performance statistics.
<code>min_targets_per_frame</code>	1	Minimum number of accepted targets required for export.
<code>min_target_score</code>	0.45	Minimum detector score for a target candidate.

Parameter	Default / configured value	Purpose
require_text	True	Requires accepted OCR output before export.
min_text_score	0.45	Minimum accepted OCR confidence.
min_export_interval_sec	0.75	Minimum time gap between accepted exports.
min_bbox_area_pixels	100.0	Minimum target bounding-box area.
min_bbox_short_side_pixels	10.0	Minimum short side of the target box.
min_image_brightness	35.0	Lower image brightness gate.
max_image_brightness	235.0	Upper image brightness gate.
max_saturated_ratio	0.25	Maximum saturated-pixel ratio.
enable_sharpness_gate	true	Enables whole-image sharpness filtering.
min_laplacian_var	120.0	Minimum whole-image Laplacian variance.
enable_roi_sharpness_gate	true	Enables target-ROI sharpness filtering.
min_roi_laplacian_var	40.0	Minimum ROI Laplacian variance.
enable_stability_gate	true	Enables temporal stability filtering across frames.
stability_required_frames	3	Number of supporting frames required for export.
stability_iou_threshold	0.40	IoU threshold for matching targets across frames.
stability_depth_tolerance_m	0.20	Depth consistency threshold in meters.
stability_track_ttl_sec	1.2	Lifetime of a temporary stability track.
stability_require_ocr_consistency	true	Requires stable OCR content across the track.
enable_dedup	true	Enables duplicate-frame rejection.
dedup_hamming_threshold	8	Hamming threshold for duplicate detection.
dedup_history_size	50	Number of recent accepted samples used for deduplication.
analysis_ocr_score_threshold	0.60	Threshold used for analysis-only OCR statistics.
distance_bin_size_m	0.25	Bin size for distance-distribution statistics.
max_analysis_distance_m	5.0	Maximum distance included in exported analysis summaries.

E.4 Detailed runtime timing breakdown

For readers interested in more detailed node-level runtime information beyond Table 4.13, Table E.4 reports the stage-level timing ranges measured for the final live-camera runtime config-

uration with OCR capped to four s.

The stage labels in the table correspond to the following execution boundaries:

- **cv_bridge**: conversion from the incoming image message to an OpenCV image.
- **pre**: preprocessing before the TensorRT engine call.
- **infer**: measured TensorRT inference time.
- **depth**: the mask-guided depth and representative-center estimation block in the YOLO node.
- **loop**: the per-target processing loop in the YOLO node.
- **post**: postprocessing after network output decoding and before publication.
- **publish**: publication of the output message.
- **total**: complete node callback time for that processing cycle.

Table E.4: Detailed stage-level timing ranges for the final live-camera runtime configuration with OCR capped to four s.

Stage	YOLO [ms]	OCR [ms]
cv_bridge	0.6–1.0	0.8–1.1
pre	0.3	1.3–1.6
infer	40.8–43.5	25.6–26.7
depth	5.3–7.0	—
loop	19.8–23.6	—
post	27.9–32.6	0.4–0.5
publish	1.2–1.3	0.5–0.7
total	70.4–77.8	33.0–35.2
Processed ROIs	—	4.0

E.5 Example final pipeline output

This section provides one concrete example of the final deployed pipeline output. The image corresponds to the university replay experiment used in Chapter 4. Figure E.1 shows the same frame rendered in two ways: a bounding-box and label view, and a center-and-distance view with the predicted segmentation masks. The abridged `/button_targets_3d_ocr` message for the same frame timestamp is shown directly below the figure.

Abridged `/button_targets_3d_ocr` message for the frame shown in Figure E.1.

```
header_stamp_ns: 1773668101805465000
frame_id: zed_left_camera_frame_optical
num_targets: 7

targets:
  - label: "ALARM"
```



Figure E.1: Example final pipeline output for frame timestamp 1773668101805465000. Left: bounding boxes with recognized button labels and detector confidence. Right: predicted masks, representative button centers, and depth estimates.

```

yolo_score: 0.893
ocr_score: 0.958
depth_m: 0.340
center_3d_m: [-0.045, 0.005, 0.340]

- label: "OPEN"
  yolo_score: 0.890
  ocr_score: 0.959
  depth_m: 0.320
  center_3d_m: [0.036, -0.001, 0.320]

- label: "CLOSE"
  yolo_score: 0.884
  ocr_score: 0.952
  depth_m: 0.332
  center_3d_m: [-0.006, 0.004, 0.332]

- label: "2"
  yolo_score: 0.863
  ocr_score: 1.000
  depth_m: 0.327
  center_3d_m: [-0.003, -0.129, 0.327]

- label: "1"
  yolo_score: 0.863
  ocr_score: 0.999
  depth_m: 0.331
  center_3d_m: [-0.005, -0.089, 0.331]

- label: "0"

```

```
yolo_score: 0.862
ocr_score: 0.997
depth_m: 0.324
center_3d_m: [-0.004, -0.045, 0.324]

- label: "3"
  yolo_score: 0.854
  ocr_score: 1.000
  depth_m: 0.327
  center_3d_m: [-0.003, -0.170, 0.327]
```